



# IoT FINAL PROJECT

## REPORT

GUO Xiaofan

YIN Chenghao

# Contents

1.	Introduction.....	- 3 -
1.1	GAP8 Introduction.....	- 3 -
1.1.1	Autonomous Peripherals .....	- 3 -
1.1.2	Ultra-low Power Microcontroller.....	- 3 -
1.1.3	Compute Engine.....	- 4 -
1.1.4	Features of GAP8.....	- 4 -
1.2	The Neural Network.....	- 4 -
1.2.1	Input Layer.....	- 6 -
1.2.2	Hidden Layer.....	- 6 -
1.2.3	Output Layer .....	- 6 -
2.	Realization .....	- 7 -
2.1	Code Link.....	- 7 -
2.2	Interpretation of Results.....	- 7 -
2.2.1	Realize a Simple Neural Network.....	- 7 -
2.2.2	Implement the Neural Network on GAP8.....	- 8 -
2.2.3	Use flex sensor data as training and test data.....	- 11 -
2.2.4	Multi-core asynchronous control .....	- 14 -
3.	Conclusions.....	- 19 -
4.	References.....	- 19 -

# 1. Introduction

In this final IoT project, the objective is to gain proficiency in utilizing the multi-core capabilities of low-power development boards GAP8, to understand the basic structure of neural networks, and to learn how to analyze performance.

The project involves mastering control of the GAP8's multi-core functionality by running different functions on distinct cores, implementing a simple neural network, and deploying this network on the GAP8. The project will use simulated data for training and predictions and will conduct a detailed performance analysis of the GAP8, focusing on aspects such as cycles, frequency, and time, as outlined in the attached document.

## 1.1 GAP8 Introduction

The GAP8 microprocessor, crafted by GreenWaves for edge computing and IoT, features a nona-core 32-bit RISC-V architecture. This innovative processor is structured around three key components: autonomous peripherals, an ultra-low power microcontroller, and the compute engine<sup>[1]</sup>.

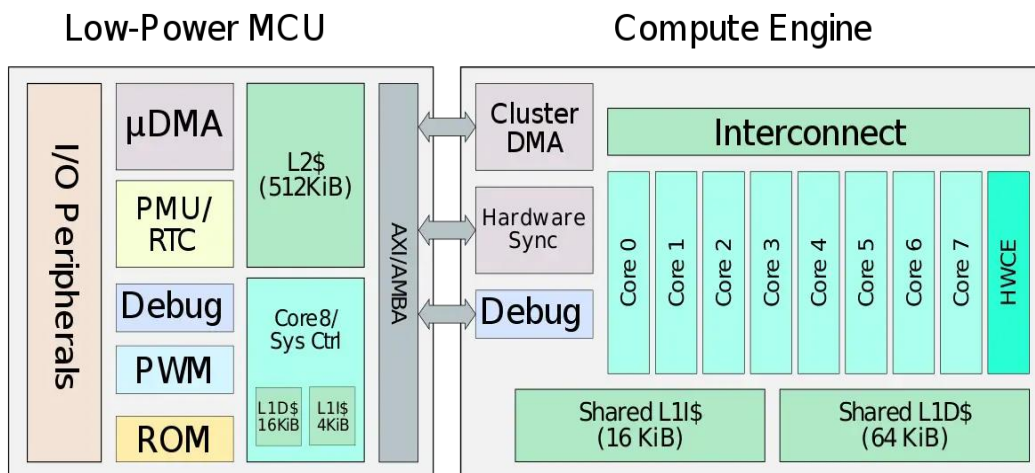


Figure 1 Gap8 Block Diagram<sup>[2]</sup>

### 1.1.1 Autonomous Peripherals

Autonomous Peripherals are hardware components that can operate independently of the main CPU. They manage data transfers between peripherals and memory without the need for CPU intervention, enabling efficient processing and low-power operation.

### 1.1.2 Ultra-low Power Microcontroller

The microcontroller block is a standard MCU (Microcontroller Unit) with many

of the standard features. The MCU is situated on its own power domain with the peripherals power switchable and configurable along with the clock generator.

The chip features nine cores (1 serving in the MCU + 8 in the compute engine), with support for the integer multiplication and division instructions (M) and compressed instructions (C) standard extensions.

The GAP8 a private L1 cache for the MCU core which consists of a 16 KiB of data cache and 4 KiB of instruction cache. The compute engine has a shared level 1 cache of its own which consists of a 16 KiB instruction cache and a 64 KiB data cache. Additionally, the entire chip shares a 512 KiB level 2 cache consisting of 4 128 KiB cache banks.

### **1.1.3 Compute Engine**

The compute engine consists of eight additional cores clustered together to form a low power but powerful computational engine. The engine sits on an entirely separate voltage and frequency domains which can be switched off when not operating or downclocked to suit a particular workload more efficiently.

### **1.1.4 Features of GAP8**

Gap8 is focused on efficient energy management, making it suitable for battery-powered or remotely deployed devices where energy is limited. Moreover, GAP8's autonomous operation allows for reduced operational costs, as it can process and analyze data independently without constant communication with a central server.

## **1.2 The Neural Network**

Neural networks are comprised of a node layers, containing an input layer, one or more hidden layers, and an output layer, rely on training data to learn and improve their accuracy over time<sup>[3]</sup>.

- Feedforward network with a single layer of neurons:

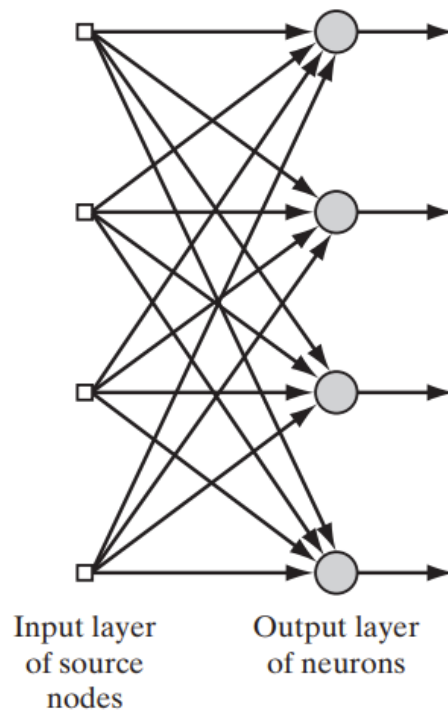


Figure 2 Feedforward network with a single layer of neurons<sup>[4]</sup>.

- Fully connected feedforward network with one hidden layer and one output layer:

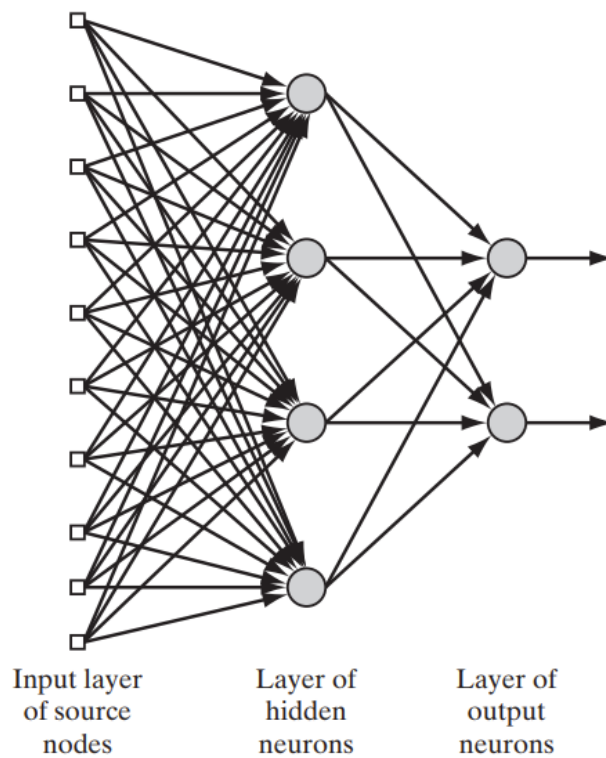


Figure 3 Fully connected feedforward network<sup>[5]</sup>.

### 1.2.1 Input Layer

The input layer is the first layer of the neural network and is responsible for receiving input data.

- **Function:** Directly process raw input data.
- **Data transformation:** The input layer usually does not make any changes to the data and passes it directly to the next layer.

### 1.2.2 Hidden Layer

Hidden layers are located between the input layer and the output layer. They are where the neural network performs calculations and feature extraction. There can be one or more hidden layers.

- **Function:** Processes data and extracts features.
- **Operation:** Each hidden layer contains multiple neurons, and each neuron is connected to all nodes from the previous layer with weights. Each neuron receives inputs from the nodes of the previous layer, applies weights, and produces an output through an activation function. This output becomes the input for the next layer. This process repeats, allowing the network to learn representations of the input data by adjusting the weights.
- **Weights:** Each connection has a weight, which represents the strength of the connection. During training, these weights are optimized according to the algorithm so that the network can perform tasks accurately.
- **Activation Function:** Introduces non-linearity to the network, enabling it to learn more complex patterns. Common activation functions include sigmoid, tanh, and ReLU.

### 1.2.3 Output Layer

The output layer is the last layer of the neural network and is responsible for outputting calculation results.

- **Function:** Provides the final output of the network, representing the solution to the problem.
- **Operation:** Each output node corresponds to a possible output category or value.

## 2.Realization

### 2.1 Code Link

- Realize a simple neural network.  
<https://github.com/heyPetiteF/IOT/blob/main/helloworld/FP1.c>
- Implement the simple neural network on GAP8.  
<https://github.com/heyPetiteF/IOT/blob/main/helloworld/FP2.c>
- Implement the neural network on GAP8 with Performance counters.  
<https://github.com/heyPetiteF/IOT/blob/main/helloworld/FP2-monitor.c>
- flexSensorData Test.  
<https://github.com/heyPetiteF/IOT/blob/main/helloworld/FP3.c>
- flexSensorData Test with Hidden\_Layer.  
<https://github.com/heyPetiteF/IOT/blob/main/helloworld/FP3-2.c>
- Multi-core asynchronous control  
<https://github.com/heyPetiteF/IOT/blob/main/helloworld/FP4.c>

### 2.2 Interpretation of Results

#### 2.2.1 Realize a Simple Neural Network

A simple neural network example implemented in C language was run and the following results were obtained:

```
*** Neural Network Test ***
Input: [0.100000, 0.120000], Target: 0.000000, Prediction: 0.011752
Input: [1.100000, 0.900000], Target: 1.000000, Prediction: 0.968554
Input: [2.100000, 1.980000], Target: 2.000000, Prediction: 2.028473
Input: [2.890000, 3.200000], Target: 3.000000, Prediction: 2.991888
```

Figure 4 The Output of a Simple NN (1).

```
*** Neural Network Test ***
Input: [0.500000, 0.400000], Target: 0.000000, Prediction: 0.093871
Input: [1.200000, 1.000000], Target: 1.000000, Prediction: 0.804686
Input: [2.500000, 2.200000], Target: 2.000000, Prediction: 2.214163
Input: [3.000000, 2.800000], Target: 3.000000, Prediction: 2.900671
```

Figure 5 The Output of a Simple NN (2).

Looking at the overall output results, the predicted value of the neural network is quite close to the target value, which indicates that the training process is successful, and the network is able to learn the relationship between the input data and the output data.

## 2.2.2 Implement the Neural Network on GAP8

### 1) Implement the Simple Neural Network on GAP8

Running a simple neural network on GAP8<sup>[6]</sup> gives the following results:

```
*** Neural Network Parameters ***

>>>Define:
Input Size: 2
Output Size: 1
Learning Rate: 0.010000
Epochs: 10000

>>>Neural network parameters:
Input: [0.100000, 0.120000], Prediction: 0.012210, Weight:0.685931
Input: [1.100000, 0.900000], Prediction: 0.968171, Weight:0.346192
Input: [2.100000, 1.980000], Prediction: 2.027989, Weight:0.000000
Input: [2.890000, 3.200000], Prediction: 2.992229, Weight:0.000000
bias:-0.097926

*** Neural network based on GAP8 ***

>>>Entering main controller
[32 0] Neural network based on GAP8
Perf : 20566
cycles Timer : 28223 cycles

>>>Cluster master core entry
[0 0] Neural network based on GAP8
[0 2] Neural network based on GAP8
[0 6] Neural network based on GAP8
[0 4] Neural network based on GAP8
[0 5] Neural network based on GAP8
[0 3] Neural network based on GAP8
[0 1] Neural network based on GAP8
[0 7] Neural network based on GAP8

>>>Cluster master core exit
Test success !
```

Figure 6 The Output of a Simple Neural Network on GAP8.

Analysis of these results shows that the neural network can produce predictions close to the target value based on the training data, indicating that the learning process is effective.

However, some values of 0.0 for weights may mean that some inputs have very little impact on the output or have been ignored by the network, which can be due to some weights contributing less to error reduction during the optimization process of the network. The final deviation value is negative, which has a small adjustment effect on the final forecast value.



Additionally, "Performance Cycles" and "Timer Cycles" provide information about the efficiency of program execution. A lower number of cycles usually means higher performance.

Finally, successful tests show that the implemented training strategy can effectively implement the forward propagation and back propagation algorithms of neural networks on GAP8 hardware.

## 2) Implement the NN on GAP8 with Performance Counters

In order to realize the function of performance counter, based on the code of "Implement the simple neural network on GAP8.", added the following three functions and run these in main function.

```

108 //uint32_t start_time;
109 clock_t start_time, end_time;
110 double total_time_in_seconds;
111
112 void start_performance_monitoring() {
113     pi_perf_conf(
114         1 << PI_PERF_CYCLES |
115         1 << PI_PERF_ACTIVE_CYCLES |
116         1 << PI_PERF_INSTR |
117         1 << PI_PERF_LD_STALL |
118         1 << PI_PERF_JR_STALL |
119         1 << PI_PERF_IMISS |
120         1 << PI_PERF_LD |
121         1 << PI_PERF_ST |
122         1 << PI_PERF_JUMP |
123         1 << PI_PERF_BRANCH |
124         1 << PI_PERF_BTAKEN |
125         1 << PI_PERF_RVC |
126         1 << PI_PERF_LD_EXT |
127         1 << PI_PERF_ST_EXT |
128         1 << PI_PERF_LD_EXT_CYC |
129         1 << PI_PERF_ST_EXT_CYC |
130         1 << PI_PERF_TCDM_CONT
131     );
132     pi_perf_reset();
133     pi_perf_start();
134     start_time = clock();
135 }
136
137 void stop_performance_monitoring() {
138     pi_perf_stop();
139     end_time = clock();
140     total_time_in_seconds = (double)(end_time - start_time) / CLOCKS_PER_SEC;
141 }
142
143 //void print_performance_data(uint32_t start_time, uint32_t end_time, uint32_t total_time) {
144 void print_performance_data() {
145     printf("Total cycles: %d\n", pi_perf_read(PI_PERF_CYCLES));
146     printf("Active cycles: %d\n", pi_perf_read(PI_PERF_ACTIVE_CYCLES));
147     printf("Instructions executed: %d\n", pi_perf_read(PI_PERF_INSTR));
148     printf("Load data stalls: %d\n", pi_perf_read(PI_PERF_LD_STALL));
149     printf("Jump register stalls: %d\n", pi_perf_read(PI_PERF_JR_STALL));
150     printf("Instruction misses: %d\n", pi_perf_read(PI_PERF_IMISS));
151     printf("Data memory loads: %d\n", pi_perf_read(PI_PERF_LD));
152     printf("Data memory stores: %d\n", pi_perf_read(PI_PERF_ST));
153     printf("Unconditional jumps: %d\n", pi_perf_read(PI_PERF_JUMP));
154     printf("Branches: %d\n", pi_perf_read(PI_PERF_BRANCH));
155     printf("Taken branches: %d\n", pi_perf_read(PI_PERF_BTAKEN));
156     printf("Compressed instructions: %d\n", pi_perf_read(PI_PERF_RVC));
157     printf("External memory loads: %d\n", pi_perf_read(PI_PERF_LD_EXT));
158     printf("External memory stores: %d\n", pi_perf_read(PI_PERF_ST_EXT));
159     printf("External load cycles: %d\n", pi_perf_read(PI_PERF_LD_EXT_CYC));
160     printf("External store cycles: %d\n", pi_perf_read(PI_PERF_ST_EXT_CYC));
161     printf("TCDM contention cycles: %d\n", pi_perf_read(PI_PERF_TCDM_CONT));
162     printf("Execution time: %f seconds\n", total_time_in_seconds);
163 }
164

```

Figure 7 Three Functions of Performance Counter.

```

166 int main(void) {
167
168     initialize();
169
170     start_performance_monitoring();
171
172     double training_data[][INPUT_SIZE] = {{0.1, 0.12}, {1.1, 0.9}, {2.1,1.98}, {2.89, 3.2}};
173
174     double targets[] = {0, 1, 2, 3};
175
176     for (int epoch = 0; epoch < EPOCHS; epoch++) {
177
178         for (int i = 0; i < sizeof(training_data) / sizeof(training_data[0]); i++) {
179
180             train(training_data[i], targets[i]);
181
182         }
183
184     }
185
186     //void stop_performance_monitoring();
187     //uint32_t end_time, total_time;
188     stop_performance_monitoring();
189
190     printf("\n\t*** Neural Network Parameters ***\n");
191
192     printf("\n>>>Define:\n");
193     printf("Input Size: %d\n", INPUT_SIZE);
194     printf("Output Size: %d\n", OUTPUT_SIZE);
195     printf("Learning Rate: %lf\n", LEARNING_RATE);
196     printf("Epochs: %d\n", EPOCHS);
197
198     printf("\n>>>Neural network parameters:\n");
199
200     //for (int i = 0; i < sizeof(training_data) / sizeof(training_data[0]); i++) {
201     for (size_t i = 0; i < sizeof(training_data) / sizeof(training_data[0]); i++) {
202
203         double prediction = predict(training_data[i]);
204
205         printf("Input: [%lf, %lf], Prediction: %lf, Weight:%lf\n", training_data[i][0],
206             training_data[i][1], prediction, weights[i]);
207     }
208
209     printf("bias:%lf\n", bias);
210
211
212     printf("\n\t *** Neural network based on GAP8 ***\n");
213
214     print_performance_data();
215
216     return pmsis_kickoff((void *)helloworld);
217 }
218 }

```

Figure 8 Called in Main Function.

However, during the running process, encountered the problem that the time function cannot be called. To capture the data of time tried using the following two functions:

1) Using function of PMSIS “pi\_time\_us()”

```

/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/BUILD
/GAP8_V3/GCC_RISCV_FREERTOS/FP2-monitor.o: In function `start_performance_monito
ring':
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/FP2-m
onitor.c:139: undefined reference to `pi_time_us'
collect2: error: ld returned 1 exit status
make: *** [/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/rtos/freertos/vendors/gwt
/rules/freertos_rules.mk:419: /home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/exampl
es/gap8/basic/helloworld/BUILD/GAP8_V3/GCC_RISCV_FREERTOS/test] Error 1

```

Figure 9 Unable to Call the Function pi\_time\_us().

## 2) Using function of C language “clock ()”

```
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/BUILD
/GAP8_V3/GCC_RISCV_FREERTOS/FP2-monitor.o: In function `start_performance_monito
ring':
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/rtos/freertos/vendors/gwt/gap8/pmsis
/include/cores/TARGET_RISCV_32/core_utils.h:15: undefined reference to `clock'
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/rtos/freertos/vendors/gwt/gap8/pmsis
/include/cores/TARGET_RISCV_32/core_utils.h:15: undefined reference to `clock'
collect2: error: ld returned 1 exit status
make: *** [/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/rtos/freertos/vendors/gwt
/rules/freertos_rules.mk:419: /home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/exampl
es/gap8/basic/helloworld/BUILD/GAP8_V3/GCC_RISCV_FREERTOS/test] Error 1
```

Figure 10 Unable to Call the Function clock().

The code compiles successfully but the result is not outputted due to the target function not being invoked.

### 2.2.3 Use flex sensor data as training and test data

#### 1) flexSensorData Test.

The following procedures were implemented for training the neural network on GAP8 using the data in the flexSensorData file:

- The flexSensorData.csv was opened with fopen().
- Each line of the file was read using fscanf().
- The parsed data were stored in the training\_data and targets arrays.
- The file was closed with fclose() upon completion of the data import.

```
/*
double training_data[100][INPUT_SIZE];
double targets[100];
int rows = 0;

#ifdef ENABLE_FILE_IO

FILE *file;
file = fopen("flexSensorData.csv", "r");
if (!file) {
    printf("CAN NOT OPEN ! \n");
    return 1;
}

while (fscanf(file, "%lf,%lf,%lf\n", &training_data[rows][0], &training_data[rows][1],
&targets[rows]) != EOF) {
    rows++;
}

fclose(file);
*/
```

Figure 11 Code to Import .csv Data Using Function.

However, a problem occurred during the reading process, the function could not be called accurately, terminal displayed as follows:

```

/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/a.c:18: undefined reference to `fopen'
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/a.c:18: undefined reference to `perror'
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/a.c:18: undefined reference to `fgets'
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/a.c:18: undefined reference to `strtok'
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/a.c:18: undefined reference to `atof'
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/a.c:18: undefined reference to `atof'
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/a.c:18: undefined reference to `strtok'
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/a.c:18: undefined reference to `fgets'
/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/a.c:18: undefined reference to `fclose'
collect2: error: ld returned 1 exit status
make: *** [/home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/rtos/freertos/vendors/gwt/rules/freertos_rules.mk:419: /home/hx/gap_riscv_toolchain_ubuntu/gap_sdk/examples/gap8/basic/helloworld/BUILD/GAP8_V3/GCC_RISCV_FREERTOS/test] Error 1

```

Figure 12 Cannot be Called from Library Function.

Therefore, modifications are made to directly introduce the data into the neural network calculation, the following results were obtained:

```

*** Neural Network Parameters ***

>>>Define:
Learning Rate: 0.000100
Epochs: 10000

>>>Neural network parameters:
Prediction: 1.412866
Weight: -0.006292
Bias: 1.979180

*** flexSensorData Basic Test ***

>>>Entering main controller
[32 0] flexSensorData Basic Test
Perf : 20095
cycles Timer : 27852 cycles

>>>Cluster master core entry
[0 0] flexSensorData Basic Test
[0 6] flexSensorData Basic Test
[0 2] flexSensorData Basic Test
[0 5] flexSensorData Basic Test
[0 7] flexSensorData Basic Test
[0 1] flexSensorData Basic Test
[0 4] flexSensorData Basic Test
[0 3] flexSensorData Basic Test

>>>Cluster master core exit
Test success !

```

Figure 13 NN Data from flexSensorData.csv (Epochs: 10000).

Considering that the weights are less than zero, the bias is greater than one, and the number of training epochs is set to a relatively high value of 10,000, there may be overfitting in the program.

Therefore, the epochs value is reduced to 100, and the results are as follows:

```
*** Neural Network Parameters ***

>>>Define:
Learning Rate: 0.000100
Epochs: 100

>>>Neural network parameters:
Prediction: 1.481679
Weight: 0.012290
Bias: 0.375596

*** flexSensorData Basic Test ***

>>>Entering main controller
[32 0] flexSensorData Basic Test
Perf : 20169
cycles Timer : 28149 cycles

>>>Cluster master core entry
[0 0] flexSensorData Basic Test
[0 4] flexSensorData Basic Test
[0 5] flexSensorData Basic Test
[0 2] flexSensorData Basic Test
[0 3] flexSensorData Basic Test
[0 6] flexSensorData Basic Test
[0 7] flexSensorData Basic Test
[0 1] flexSensorData Basic Test

>>>Cluster master core exit
Test success !
```

Figure 14 NN Data From flexSensorData.csv (Epochs:100).

The specific values of the weights and biases of the neural network after a certain number of trainings provide some information about the state of the network after training. Performance indicators indicate that the program can run on the cluster and the execution time is reasonable.

## 2) flexSensorData Test with Hidden Layer

Considering that neural networks with hidden layers have stronger learning capabilities and can theoretically approximate any continuous function, we made the following attempt.

Based on the data in the flexSensorData.csv file, a neural network with hidden layers was trained on GAP8 and the following results are obtained:

```

*** Neural Network Parameters ***

>>>Define:
Learning Rate: 0.010000
Epochs: 10000
HIDDEN_LAYER_SIZE: 10

>>>Neural network parameters:
Weight: 1.707808
Prediction: 1.686064
Bias: 1.686064

*** flexSensorData HIDDEN_LAYER Test ***

>>>Entering main controller
[32 0] flexSensorData HIDDEN_LAYER Test
Perf : 20399
cycles Timer : 28733 cycles

>>>Cluster master core entry
[0 0] flexSensorData HIDDEN_LAYER Test
[0 4] flexSensorData HIDDEN_LAYER Test
[0 5] flexSensorData HIDDEN_LAYER Test
[0 7] flexSensorData HIDDEN_LAYER Test
[0 6] flexSensorData HIDDEN_LAYER Test
[0 1] flexSensorData HIDDEN_LAYER Test
[0 3] flexSensorData HIDDEN_LAYER Test
[0 2] flexSensorData HIDDEN_LAYER Test

>>>Cluster master core exit
Test success !

```

Figure 15 NN with Hidde\_Layer.

## 2.2.4 Multi-core asynchronous control

```

79 void function_core_0(void *arg) {
80
81     initialize();
82
83     double training_data[][INPUT_SIZE] = {{0.1, 0.12}, {1.1, 0.9}, {2.1, 1.98}, {2.89, 3.2}};
84     double targets[] = {0, 1, 2, 3};
85
86     for (int epoch = 0; epoch < EPOCHS; epoch++) {
87         for (int i = 0; i < sizeof(training_data) / sizeof(training_data[0]); i++) {
88             train(training_data[i], targets[i]);
89         }
90     }
91     printf("\n\t***[Core 0] Training the simple NN on the GAP8***");
92     for (int i = 0; i < sizeof(training_data) / sizeof(training_data[0]); i++) {
93         double prediction = predict(training_data[i]);
94         printf("\n[Core 0]Input: [%.2f, %.2f], Target: %.2f, Prediction: %lf\n",
95             training_data[i][0], training_data[i][1], targets[i], prediction);
96     }
97 }

```

Figure 16 Simple NN Run in the 1st Core.



```

101 void function_core_1(void *arg) {
102
103     //initialize();
104
105     const int NUM_SAMPLES = 4;
106     double random_data[NUM_SAMPLES][INPUT_SIZE];
107
108     for (int i = 0; i < NUM_SAMPLES; i++) {
109         for (int j = 0; j < INPUT_SIZE; j++) {
110             random_data[i][j] = ((double)simple_rand() / RAND_MAX) * 2 - 1;
111         }
112     }
113     printf("\n\t***[Core 1] Generating random prediction dataset***");
114     for (int i = 0; i < NUM_SAMPLES; i++) {
115         double prediction = predict(random_data[i]);
116         printf("\n[Core 1]Random Input: [%lf, %lf], Prediction: %lf\n",
117             random_data[i][0], random_data[i][1], prediction);
118     }
119
120 }
121

```

Figure 17 Stochastic NN Run in the 2nd Core.

```

124 void function_core_2(void *arg) {
125
126     initialize();
127
128     double weights_input_hidden[INPUT_SIZE][HIDDEN_SIZE] = {{0.15, 0.25}, {0.20, 0.30}};
129     double weights_hidden_output[HIDDEN_SIZE][OUTPUT_SIZE] = {{0.40}, {0.50}};
130     double input[INPUT_SIZE] = {1.0, 2.0};
131     double hidden[HIDDEN_SIZE] = {0};
132     double output[OUTPUT_SIZE] = {0};
133
134     for (int i = 0; i < HIDDEN_SIZE; i++) {
135         for (int j = 0; j < INPUT_SIZE; j++) {
136             hidden[i] += input[j] * weights_input_hidden[j][i];
137         }
138         hidden[i] = activation(hidden[i]);
139     }
140
141     for (int i = 0; i < OUTPUT_SIZE; i++) {
142         for (int j = 0; j < HIDDEN_SIZE; j++) {
143             output[i] += hidden[j] * weights_hidden_output[j][i];
144         }
145         output[i] = activation(output[i]);
146     }
147     printf("\n\t***[Core 2] NN forward propagation function with Hidden_Layer***");
148     printf("\n[Core 2]Input:(%.2f,%.2f), hidden:%f, Output: %f\n", input[0],input[1],
149         hidden[0], output[0]);
150 }

```

Figure 18 NN Forward Propagation with Hidden Layer Run in the 3rd Core.

```

152 void function_core_3(void *arg) {
153
154     initialize();
155
156     double weights[2] = {0.5, -0.5};
157     double input[2] = {0.5, 0.6};
158     double expected_output = 0.7;
159     double output;
160
161     output = activation(input[0] * weights[0] + input[1] * weights[1]);
162
163     double error = expected_output - output;
164
165     for (int i = 0; i < 2; i++) {
166         double gradient = error * activation_derivative(output) * input[i];
167         weights[i] += LEARNING_RATE * gradient;
168     }
169
170     printf("\n\t***[Core 3] NN back propagation function with Hidden_Layer***");
171     printf("\n[Core 3]input:(%.2f,%.2f),Updated weights: %f, %f\n",input[0],input[1],
172     weights[0], weights[1]);
173 }

```

Figure 19 NN Back Propagation with Hidden Layer Run in the 4th Core.

```

174 void function_core_4(void *arg) {
175
176     printf("\n\t***[Core 4] Linear Regression Model***");
177
178     double weights = 0.0;
179     double bias = 0.0;
180     double inputs[] = {1.0, 2.0, 3.0, 4.0};
181     double targets[] = {2.0, 4.0, 6.0, 8.0};
182     int n_samples = sizeof(inputs) / sizeof(inputs[0]);
183
184     for (int epoch = 0; epoch < EPOCHS; epoch++) {
185         for (int i = 0; i < n_samples; i++) {
186             double output = inputs[i] * weights + bias;
187             double error = targets[i] - output;
188             weights += LEARNING_RATE * error * inputs[i];
189             bias += LEARNING_RATE * error;
190         }
191     }
192     printf("\n[Core 4]Trained weights: %f, bias: %f\n", weights, bias);
193 }
194

```

Figure 20 Linear Regression Run in the 5th Core.

```

195 void function_core_5(void *arg) {
196
197     printf("\n\t***[Core 5] Logistic Regression Model***");
198
199     double weights = 0.0;
200     double bias = 0.0;
201     double inputs[] = {0, 1, 2, 3};
202     double targets[] = {0, 0, 1, 1};
203     int n_samples = sizeof(inputs) / sizeof(inputs[0]);
204
205     for (int epoch = 0; epoch < EPOCHS; epoch++) {
206         for (int i = 0; i < n_samples; i++) {
207             double output = activation(inputs[i] * weights + bias);
208             double error = targets[i] - output;
209             weights += LEARNING_RATE * error * output * (1 - output) * inputs[i];
210             bias += LEARNING_RATE * error * output * (1 - output);
211         }
212     }
213     printf("\n[Core 5]Trained weights: %f, bias: %f\n", weights, bias);
214 }
215

```

Figure 21 Logistic Regression Run in the 6th Core.



```

216 void function_core_6(void *arg) {
217
218     printf("\n\t***[Core 6] Multi-Layer Perceptron Model***");
219
220     double weights1 = 0.15, weights2 = 0.25;
221     double bias1 = 0.35, bias2 = 0.45;
222
223     double inputs[] = {0.5, 0.6};
224     double target = 0.7;
225     double hidden_output, final_output;
226
227     for (int epoch = 0; epoch < EPOCHS; epoch++) {
228
229         hidden_output = activation(inputs[0] * weights1 + inputs[1] * weights2 + bias1);
230         final_output = activation(hidden_output * weights2 + bias2);
231
232         double error = target - final_output;
233         weights2 += LEARNING_RATE * error * final_output * (1 - final_output) * hidden_output;
234         weights1 += LEARNING_RATE * error * final_output * (1 - final_output) * inputs[0];
235         bias2 += LEARNING_RATE * error * final_output * (1 - final_output);
236         bias1 += LEARNING_RATE * error * final_output * (1 - final_output);
237     }
238
239     printf("\n[Core 6]Trained weights1: %f, weights2: %f, bias1: %f, bias2: %f\n", weights1,
240           weights2, bias1, bias2);
241 }

```

Figure 22 Multi-Layer Perceptron Run in the 7th Core.

```

242 void function_core_7(void *arg) {
243
244     printf("\n\t***[Core 7] Simple Classification Model***");
245
246     double weights[] = {0.0, 0.0};
247     double bias = 0.0;
248     double inputs[][2] = {{0, 0}, {0, 1}, {1, 0}, {1, 1}};
249     double targets[] = {0, 1, 1, 0}; // XOR-like problem
250     int n_samples = sizeof(inputs) / sizeof(inputs[0]);
251
252     for (int epoch = 0; epoch < EPOCHS; epoch++) {
253         for (int i = 0; i < n_samples; i++) {
254             double output = activation(inputs[i][0] * weights[0] + inputs[i][1] * weights[1] +
255             bias);
256             double error = targets[i] - output;
257             weights[0] += LEARNING_RATE * error * output * (1 - output) * inputs[i][0];
258             weights[1] += LEARNING_RATE * error * output * (1 - output) * inputs[i][1];
259             bias += LEARNING_RATE * error * output * (1 - output);
260         }
261     }
262     printf("\n[Core 7]Trained weights: [%f, %f], bias: %f\n", weights[0], weights[1], bias);
263 }
264

```

Figure 23 Simple Classification Run in the 8th Core.

Run the code and get the following results:

```
>>>Entering main controller
>>>Entering cluster on core 0

***[Core 3] NN back propagation function with Hidden_Layer***
***[Core 2] NN forward propagation function with Hidden_Layer***
***[Core 1] Generating random prediction dataset***
[Core 2]Input:(1.00,2.00), hidden:0.550000, Output: 0.645000
[Core 1]Random Input: [-0.258156, -0.672048], Prediction: 0.309833
[Core 3]input:(0.50,0.60),Updated weights: 0.500000, -0.500000

[Core 1]Random Input: [-0.184790, 0.921812], Prediction: 0.000000

[Core 1]Random Input: [0.916379, 0.919858], Prediction: 0.067947

[Core 1]Random Input: [-0.071200, 0.602588], Prediction: 0.196470
***[Core 6] Multi-Layer Perceptron Model***
[Core 6]Trained weights1: 0.179153, weights2: 0.286263, bias1: 0.408306, bias2: 0.508306
***[Core 4] Linear Regression Model***
[Core 4]Trained weights: 2.000000, bias: 0.000000
***[Core 5] Logistic Regression Model***
[Core 5]Trained weights: 0.000000, bias: 0.000000
***[Core 7] Simple Classification Model***
[Core 7]Trained weights: [0.000000, 0.000000], bias: 0.000000

***[Core 0] Training the simple NN on the GAP8***
[Core 0]Input: [0.10, 0.12], Target: 0.00, Prediction: 0.012123

[Core 0]Input: [1.10, 0.90], Target: 1.00, Prediction: 0.968244

[Core 0]Input: [2.10, 1.98], Target: 2.00, Prediction: 2.028082

[Core 0]Input: [2.89, 3.20], Target: 3.00, Prediction: 2.992164

>>>Leaving cluster on core 0
>>>Leaving main controller
```

Figure 24 Results of Function Run in Different Cores

It can be seen from the results that the titles and output results of different functions which running in different Cores are mixed and are not displayed in the order.

Each “function\_core\_n(void \*arg)” in the output results is executed on a different core, but their output are redirected to the same standard output stream: Terminal. Therefore, the output results are displayed continuously rather than separated by cores. In another world, their output are displayed consecutively in the order of execution, not in programming order.

### 3. Conclusions

In this final project learned about the architecture and functions of GAP8 and gained a preliminary understanding and application of neural networks.

The project mainly completed the successful deployment of a simple neural network on the GAP8 platform, the use data of flexSensorData.csv for training and testing, and the method of multi-core asynchronous control.

The goals were basically completed, but some problems were still encountered in the process, such as the linking and function calling of the C language function library and PMSIS function library on the GAP8.

### 4. References

- [1] GAP8 – GreenWaves, <https://en.wikichip.org/wiki/greenwaves/gap8>
- [2] File: gap8 block diagram.svg, [https://en.wikichip.org/wiki/File:gap8\\_block\\_diagram.svg](https://en.wikichip.org/wiki/File:gap8_block_diagram.svg)
- [3] What are neural networks? <https://www.ibm.com/topics/neural-networks>
- [4] Haykin, S. (2009). Neural networks and learning machines (3rd ed., p. 21). Upper Saddle River, NJ: Pearson Education, Inc.
- [5] Haykin, S. (2009). Neural networks and learning machines (3rd ed., p. 22). Upper Saddle River, NJ: Pearson Education, Inc.
- [6] <https://greenwaves-technologies.com/manuals/BUILD/PULP-OS/html/index.html#section1>