



# **CASSANDRA - INTRODUCTION AND PRACTICAL LAB**

*Rafael Angarita*

*ISEP - Paris*

*[rafael.angarita@isep.fr](mailto:rafael.angarita@isep.fr)*

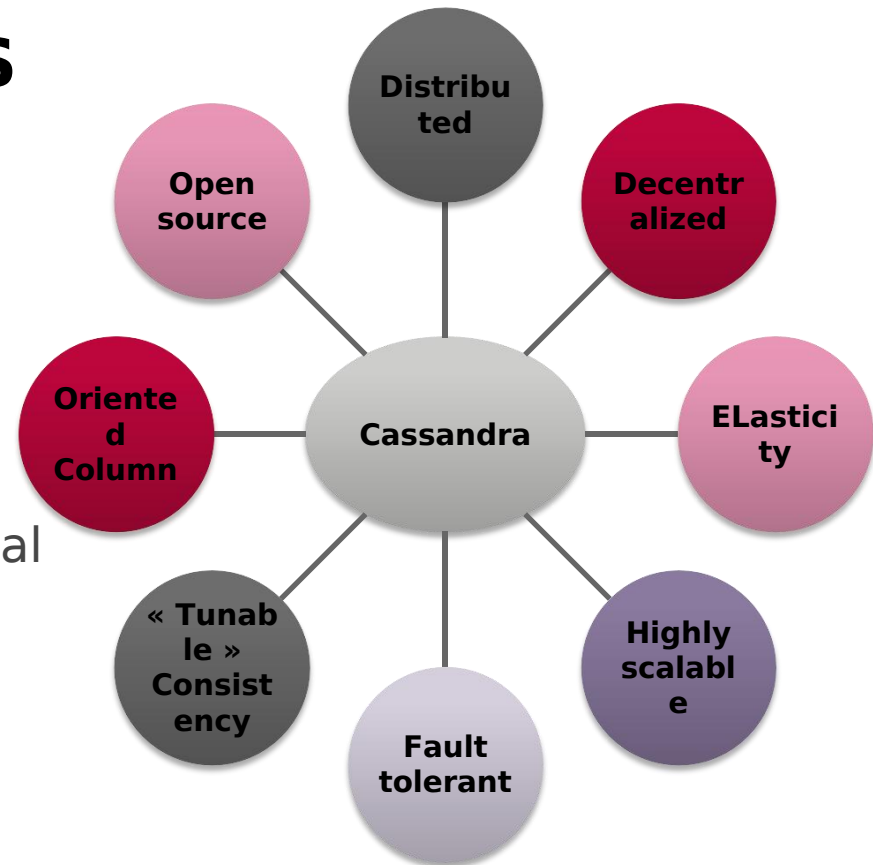
# CASSANDRA

- Initiator: Facebook in 2007 (solution for the problem "inbox search »)
- Apache 2.0 License
- Written in Java
- Data Model
  - Based on BigTable (data model) and Dynamo (partitioning and consistency)
- Thrift interface: Ruby, Perl, Python, Scala and Java, ...
- CQL (Cassandra Query Language): SQL-Like



# FONCTIONNALITÉS

- Designed to handle large amount of data across multiple servers
- Easy to implement and deploy
- Mimics traditional relational database systems



# CARACTÉRISTICS



## SUITABLE FOR SPARSE DATA

Primary Key	First Name	Last Name	E-mail ID
1	Avril	D'Souza	NULL 😞
2	David	Gomes	davidgomes1@yahoo.com
3	Susane	NULL 😞	NULL 😞

First Name	Last Name
Avril	D'Souza

First Name	Last Name	E-mail ID
David	Gomes	davidgomes1@yahoo.com


First Name
Susane

# PERFORMANCE

Its design allows it to surpass competing DBs

Very good read / write rates: Improves linearly with the addition of new nodes

*"In terms of scalability, there is a clear winner throughout our experiments. Cassandra achieves the highest throughput for the maximum number of nodes..." - University of Toronto*



**CASSANDRA**






# **CASSANDRA ARCHITECTURE**

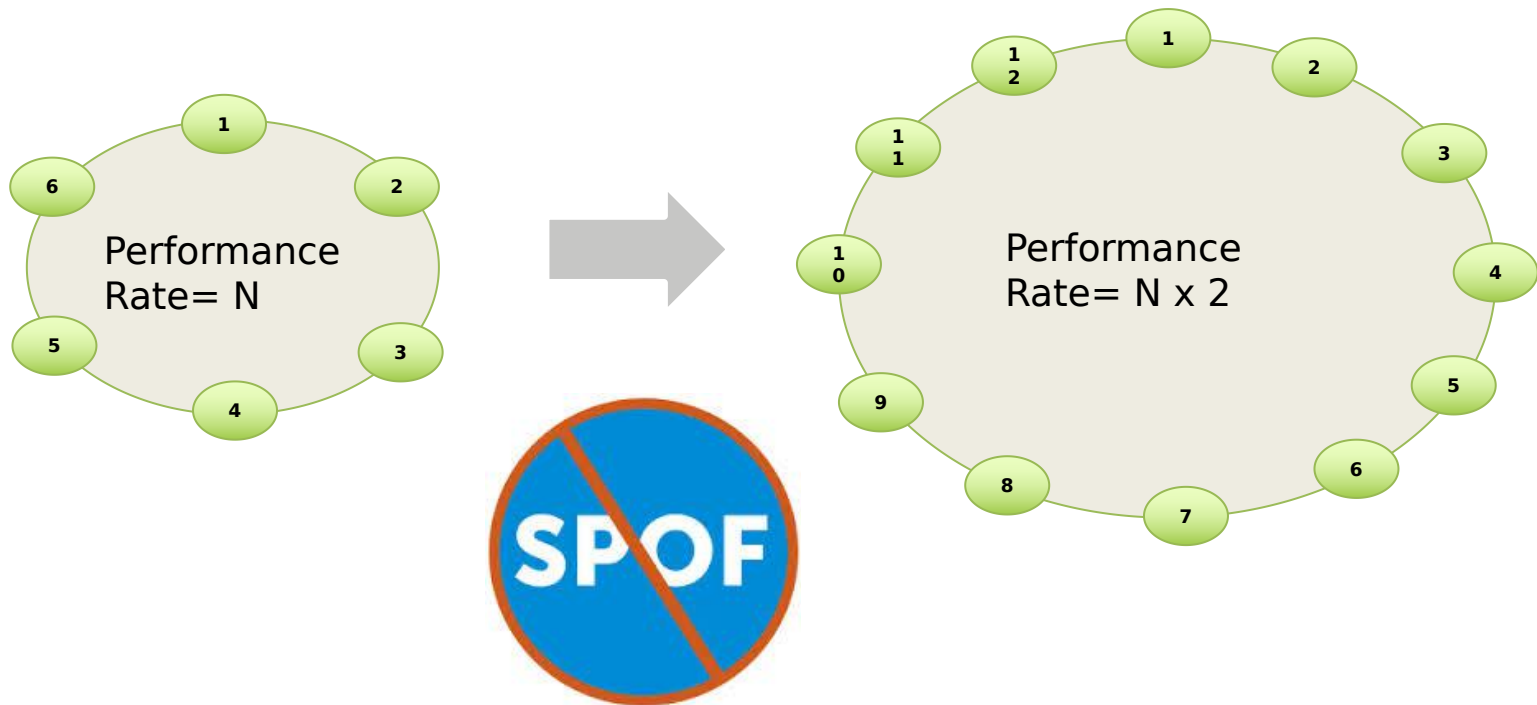


# OVERVIEW

- Cassandra was designed with the understanding that system/ hardware failures can and do occur
- Peer-to-peer, distributed system
- All nodes are the same
- Data partitioned among all nodes in the cluster
- Custom data replication to ensure fault tolerance
- Read/Write-anywhere design
- Google BigTable - data model
  - Column Families
  - Memtables
  - SSTables
- Amazon Dynamo - distributed systems technology
  - Consistent hashing
  - Partitioning
  - Replication
  - One-hop routing




# SCALING OUT ET HIGH AVAILABILITY



**CASSANDRA**

# WRITE OPERATIONS



# WRITE OPERATIONS

- Commit log
  - First place a write is recorded
  - Crash recovery mechanism
- MemTable
  - Data structure in memory
  - Once recorded in commit log, data is written to Memtable
  - Once memtable size reaches a threshold, it is flushed (appended) to SSTable
  - First place read operations look for data
- SSTable
  - Kept on disk
  - Immutable once written
  - Periodically compacted for performance

# CONSISTENCY IN CASSANDRA

- ℓ Architecture « **Read and Write Everywhere** »
- ℓ The user can connect to any node, any data center, and read / write the data he wants
- ℓ Cassandra the fastest NoSQL database in **write operations**
- ℓ Extension of the notion of eventual consistency to **a tunable consistency**
- ℓ The choice of the consistency is made by request: clause USING CONSISTENCY (the ONE level is by default)


# write consistency

Level	Description
ONE	1 replica.
QUORUM	$(N / 2) + 1$
ALL	$N = \text{replication factor}$

# read consistency

Level	Description
ONE	1 replica
QUORUM	Return most recent TS after $(N / 2) + 1$ report
ALL	$N = \text{replication factor}$


## WRITE- CONSISTENCY ONE



INSERT INTO table (column1, ...) VALUES (value1, ...) USING  
CONSISTENCY ONE



## WRITE - CONSISTENCY QUORUM



INSERT INTO table (column1, ...) VALUES (value1, ...) USING  
CONSISTENCY QUORUM

**CASSANDRA**

## GENERALITIES

- Cassandra is designed as a distributed database management system
  - use it when you have a lot of data spread across multiple servers
- Cassandra write performance is always excellent, but read performance depends on write patterns
  - it is important to spend enough time to design proper schema around the query pattern

## STRENGTHS AND WEAKNESSES

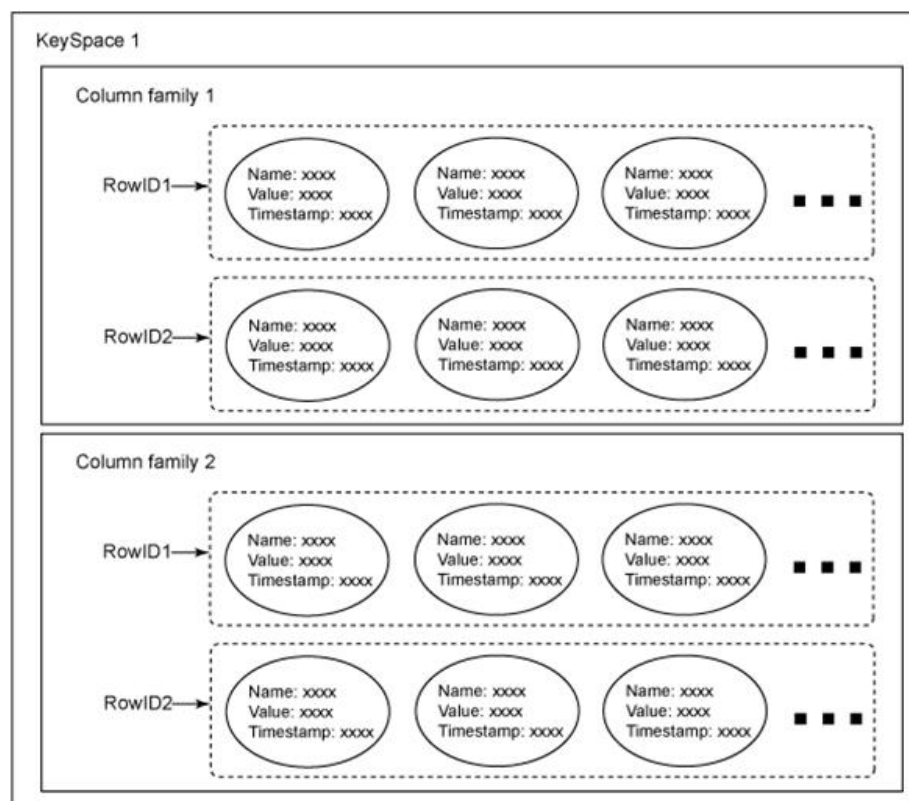
- perfect for time-series data
- high performance
- Decentralization
- nearly linear scalability
- replication support
- no single points of failure
- **MapReduce** support

- no referential integrity
- no concept of JOIN
- querying options for retrieving data are limited
- sorting data is a design decision
- no GROUP BY
- first think about queries, then about data model

# DATA MODEL

# KEY-COLUMN(S) MODEL

- A record is a collection of labeled columns (with a name)
- Family of columns = Table (By analogy with RDBMS)
- A record must contain at least one column



# EXAMPLE

Column Name

Column Family: Tweets

1234e530-8b82-11df	Text	User_ID	Date
	Hello, World!	39823	2009-03-25T19:20:30

22615e20-8b82-11df	Text	User_ID	Date
	Gooooal!	592	2009-03-25T19:25:43

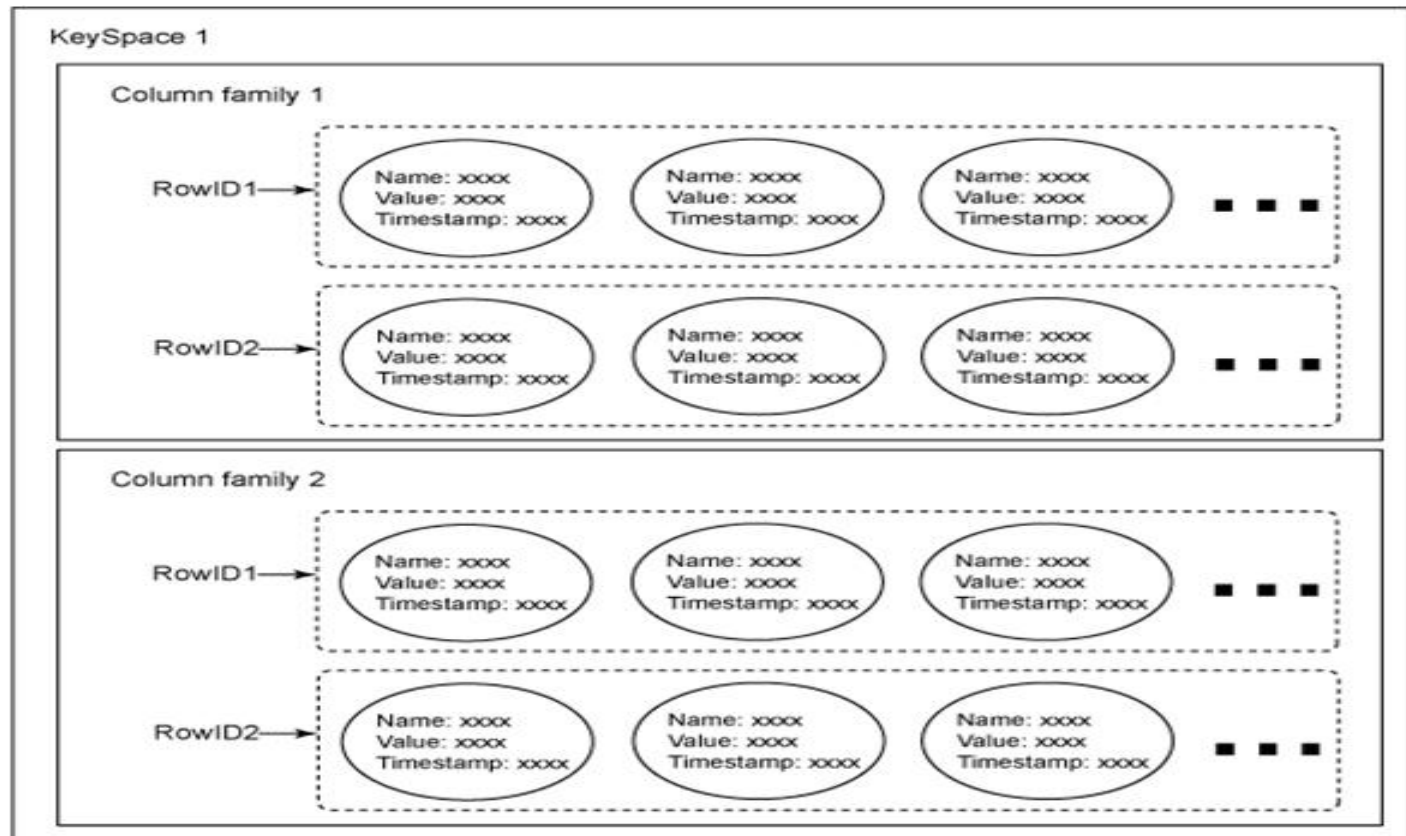
• • •

Key

Column Value

# KEYSPACE

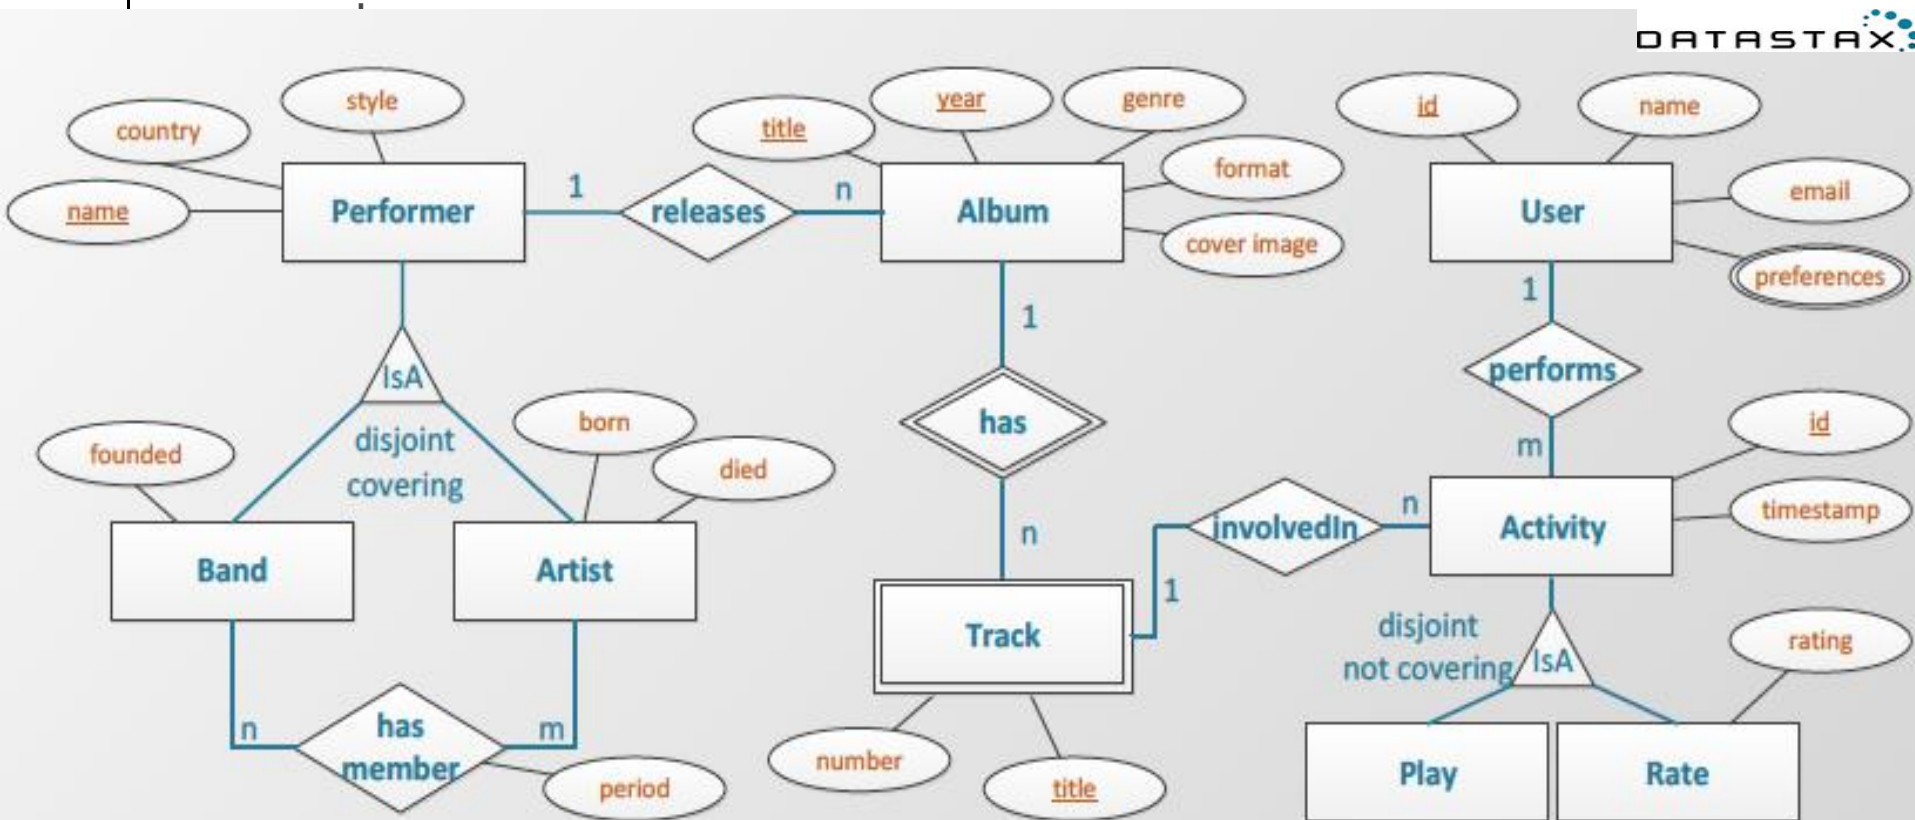
Set of column families (~ Database)



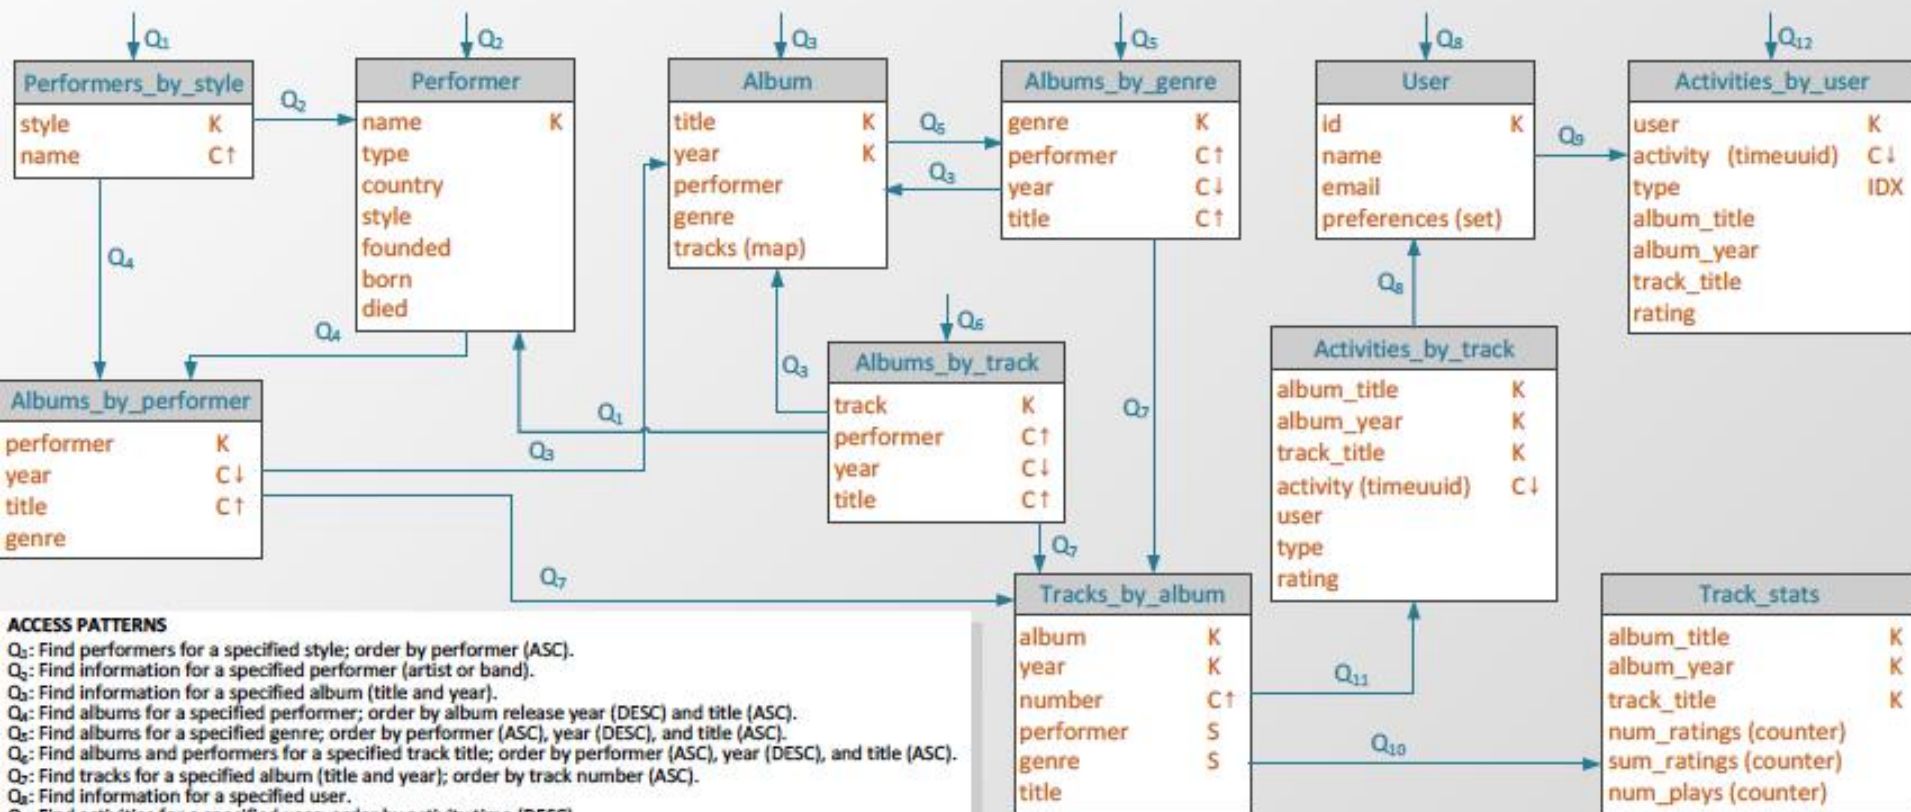


# METHODOLOGY - E/R MODEL

Diagramm ER (Chen): entities, associations, cardinalities,



# DATA MODEL FOR CASSANDRA



## ACCESS PATTERNS

- Q1: Find performers for a specified style; order by performer (ASC).
- Q2: Find information for a specified performer (artist or band).
- Q3: Find information for a specified album (title and year).
- Q4: Find albums for a specified performer; order by album release year (DESC) and title (ASC).
- Q5: Find albums for a specified genre; order by performer (ASC), year (DESC), and title (ASC).
- Q6: Find albums and performers for a specified track title; order by performer (ASC), year (DESC), and title (ASC).
- Q7: Find tracks for a specified album (title and year); order by track number (ASC).
- Q8: Find information for a specified user.
- Q9: Find activities for a specified user; order by activity time (DESC).
- Q10: Find statistics for a specified track.
- Q11: Find user activities for a specified track; order by activity time (DESC).
- Q12: Find user activities for a specified activity type.

# **CASSANDRA QUERY LANGUAGE - CQL**

# KEYSPACE

```
CREATE KEYSPACE demo
    WITH replication = {'class': 'SimpleStrategy',
        replication_factor': 3};
```

```
USE demo;
```

```
DROP KEYSPACE demo;
```

## CREATING A TABLE (COLUMN FAMILY)

```
CREATE TABLE users(
  email varchar,
  bio varchar,
  birthday timestamp,
  active boolean,
  PRIMARY KEY (email));
```

```
CREATE TABLE tweets(
  email varchar PRIMARY
  KEY,
  time_posted timestamp,
  tweet varchar);
```

## INSERTION

```
INSERT INTO users (email, bio, birthday, active)
VALUES ('isep.rdi@gmail.com', 'Associate
professor',
516513600000, true);
```

Import data from csv file

```
COPY table1 (column1, column2, column3, column4)
FROM 'data.csv';
```

With header

```
COPY table1 (column1, column2, column3, column4)
FROM 'data.csv'
```

```
WITH HEADER=true;
```

**CASSANDRA**

## QUERYING

```
SELECT * FROM users;
```

```
SELECT count(*) from users;
```

```
SELECT * FROM users LIMIT 10;
```

```
SELECT email FROM users WHERE active = true;
```



# DENORMALISATION

## REMINDERS - EXAMPLE

**videos**

id	title	runtime	year
1	Insurgent	119	2015
2	Interstellar	98	2014
3	Mockingjay	122	2014
...	...	...	...

**users**

id	login	name
a	emotions	Mr. Emotional
b	clueless	Mr. Naïve
c	noshow	Mr. Inactive
...	...	...

**comments**

id	user_id	video_id	comment
1	a	1	Loved it!
2	a	3	Hated it!
3	a	2	I cried at the end!
4	b	2	Someone stole my tissues...
...	...	...	...

## COMMENTS ON EACH VIDEO

```
SELECT comment
FROM videos JOIN comments ON videos.id =
comments.video_id
```

videos JOIN comments

id	title	runtime	year	id	user_id	video_id	comment
1	Insurgent	119	2015	1	a	1	Loved it!
3	Mockingjay	122	2014	2	a	3	Hated it!
2	Interstellar	98	2014	3	a	2	I cried at the end!
2	Interstellar	98	2014	4	b	2	Someone stole my tissues.
...	...	...	...	...	...	...	...

# COMMENTS ON EACH LOGIN AND USER

users

id	login	name
a	emotions	Mr. Emotional
b	clueless	Mr. Naïve
c	noshow	Mr. Inactive
...	...	...
...	...	...
...	...	...
...	...	...
...	...	...
...	...	...
...	...	...
...	...	...

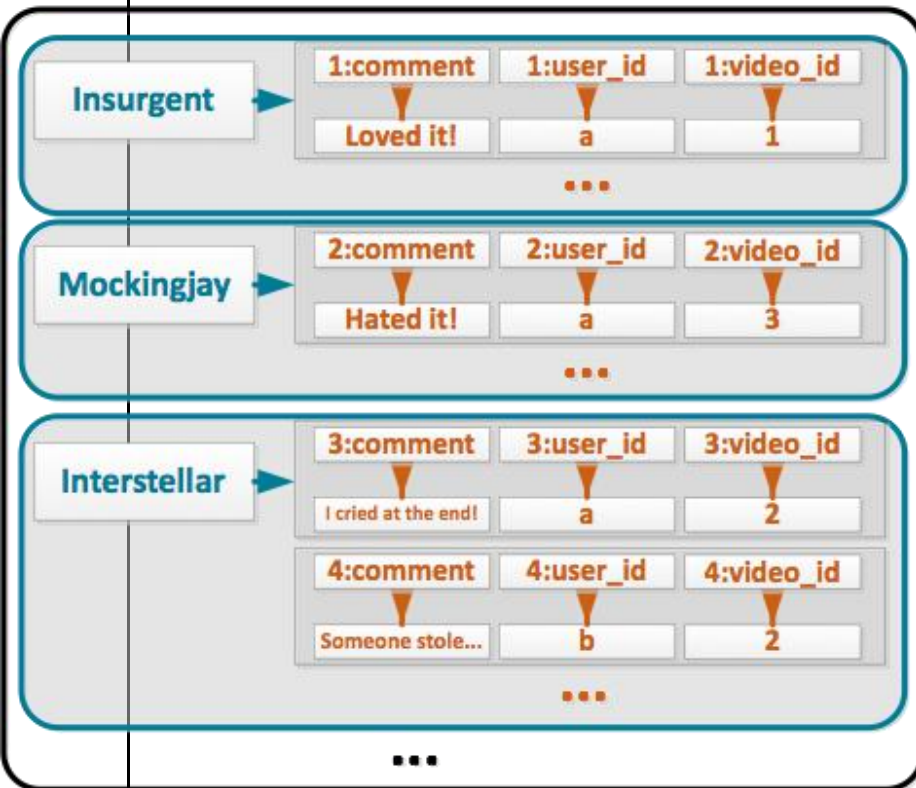
comments

id	user_id	video_id	comment
1	a	1	Loved it!
2	a	3	Hated it!
3	a	2	I cried at the end!
4	b	2	Someone stole my tissues..
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...

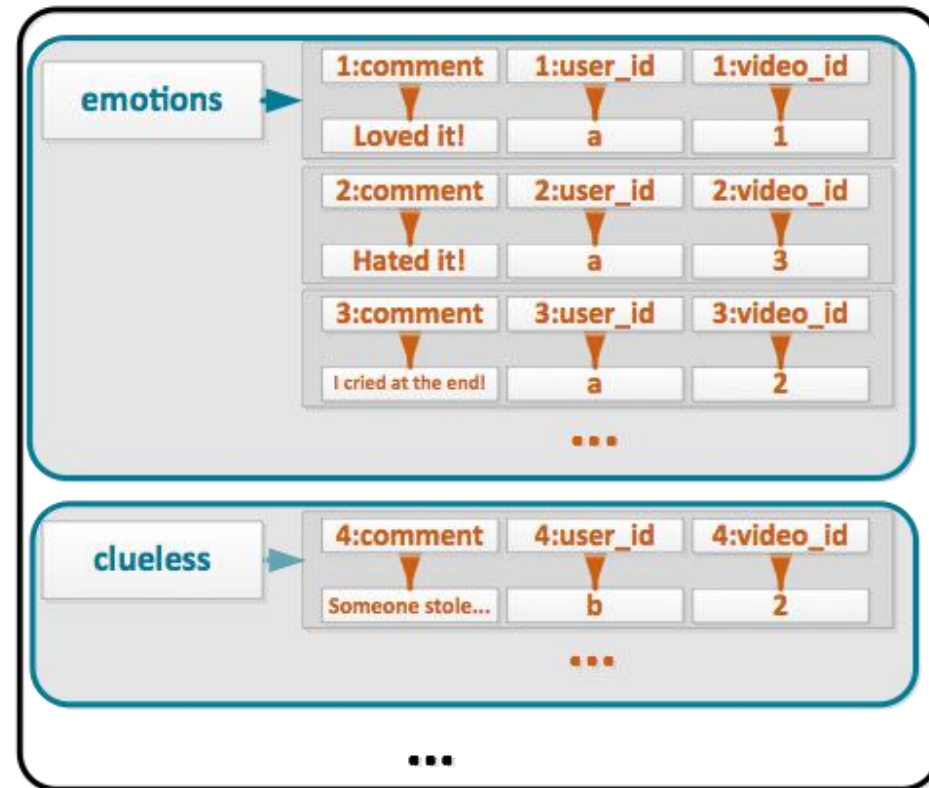


# DENORMALISATION IN CASSANDRA

comments\_by\_video



comments\_by\_user



# TABLES

Create tables and insert records

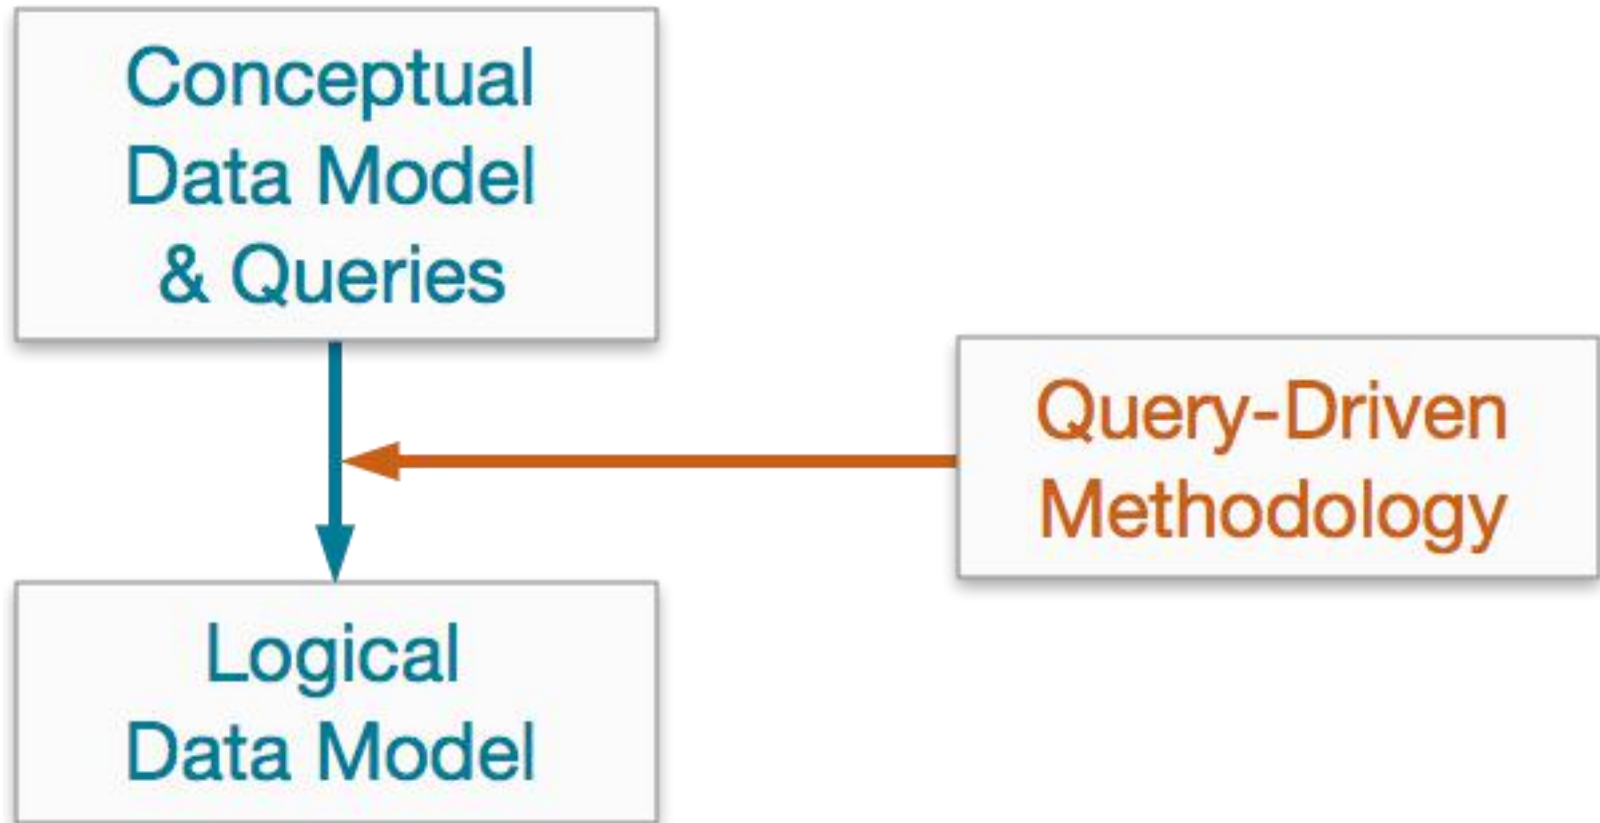
Write queries to find comments for a particular movie / user

We want to classify the videos commented by a user from the most recent to the oldest. What do we need to do?

# **CONCEPTUAL DATA MODEL (CDM)**



## TO THE DATA LOGIC MODEL



# LOGIC MODEL: CHEBOTKO DIAGRAM

Diagram of tables and queries (access patterns)

Q1 ↓

videos	
video_id	K
uploaded_timestamp	
title	
description	
type	
url	
*encoding*	
{tags}	
<preview_thumbnails>	

Q2 →

actors_by_video	
video_id	K
actor_name	C↑
character_name	C↑

## UDTs

\*encoding\*

encoding  
height  
width  
{bit\_rates}

## ACCESS PATTERNS

Q1: Find a video with a **specified video id**

Q2: Find actors for a **video with a known id**(show actor names in ascending order)

## CHEBOTKO DIAGRAM: NOTATION

table_name		
column_name_1	CQL Type K	← Partition key column
column_name_2	CQL Type C↑	← Clustering key column (ASC)
column_name_3	CQL Type C↓	← Clustering key column (DESC)
column_name_4	CQL Type S	← Static column
column_name_5	CQL Type IDX	← Secondary index column
column_name_6	CQL Type ++	← Counter column
[column_name_7]	CQL Type	← Collection column (list)
{column_name_8}	CQL Type	← Collection column (set)
<column_name_9>	CQL Type	← Collection column (map)
*column_name_10*	UDT Name	← UDT column
(column_name_11)	CQL Type	← Tuple column
column_name_12	CQL Type	← Regular column

## **CASSANDRA-CONCLUSION**

perfect for time-series data

high performance

Decentralization

nearly linear scalability

replication support

no single points of failure

MapReduce support

first think about queries, then about data model

## MORE INFORMATION

Dev: <http://www.datastax.com/dev>

Docs: <http://docs.datastax.com/en/index.html>

Planet Cassandra: <http://planetcassandra.org/>

blogs: <http://tobert.github.io/>

<http://patrickmcfadin.com/>

<http://rustyrazorblade.com/>

<https://ahappyknockoutmouse.wordpress.com/author/anukeus/>

<http://thelastpickle.com/blog/>

Livre : <http://www.amazon.com/>

Cassandra-High-Availability-Robbie-Strickland/dp/1783989122

DataStax Academy: <https://academy.datastax.com/>

Formation: <http://www.datastax.com/what-weoffer/products-services/training>