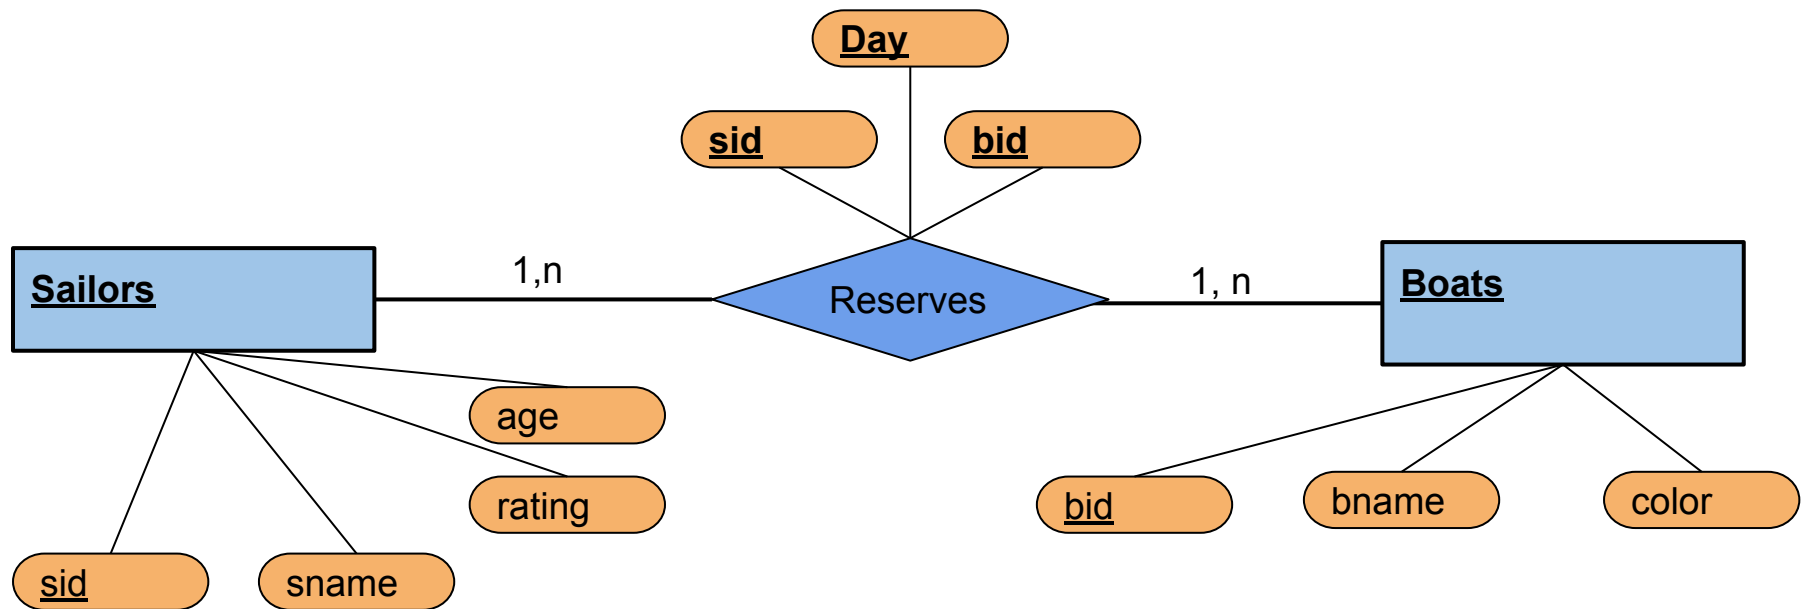


SQL: Queries, Constraints

Our Database

- ❖ Sailors(*sid: integer, sname: string, rating: integer, age: real*)
- ❖ Boats(*bid: integer, bname: string, color: string*)
- ❖ Reserves (*sid: integer, bid: integer, day: date*)

Our Model: Schema



Basic SQL Query

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>

- ❖ *target-list* A list of attributes of relations in *relation-list*
- ❖ *relation-list* A list of relation names.
- ❖ *qualification* Comparisons (Attr *op* const or Attr1 *op* Attr2, where *op* is one of $<$, $>$, $=$, \leq , \geq , \neq) combined using AND, OR and NOT.
- ❖ **DISTINCT** is an optional keyword indicating that the answer should not contain duplicates. Default is that duplicates are not eliminated!

Examples

- ❖ Produce the names of sailors:

```
SELECT sname  
FROM Sailors
```

SNAME
dustin
lubber
rusty

- ❖ Change the column heading:

```
SELECT sname as "Sailor Name"  
FROM Sailors
```

Sailor Name
dustin
lubber
rusty

- ❖ Can rename table names as well :

```
SELECT S.sname  
FROM Sailors S
```

Examples (2)

- ❖ Find names and ages of all sailors.

```
SELECT  DISTINCT S.sname, S.age  
        FROM  Sailors S;
```

```
SELECT  DISTINCT Sailors.sname, Sailors.age  
        FROM  Sailors;
```

```
SELECT  DISTINCT sname, age  
        FROM  Sailors;
```

Examples (3)

- ❖ Find sailors with rating above 7.

```
SELECT    S.sid, S.sname, S.rating, S.age
FROM Sailors S
WHERE     S.rating > 7;
```

```
SELECT    *
FROM Sailors
WHERE     rating > 7;
```

LIMIT and ORDER BY

- ❖ LIMIT: If you want to get limit the number of results
Get the 10 first Sailors

```
SELECT * FROM Sailors LIMIT 10
```

- ❖ ORDER BY: order results by an attribute

```
SELECT * FROM Sailors ORDER BY age
```

- ❖ ORDER BY and LIMIT: get TOP 10 Youngest Sailors

```
SELECT * FROM Sailors ORDER BY age LIMIT 10
```


Conceptual Evaluation Strategy

- ❖ Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
 - Compute the cross-product of *relation-list*.
 - Discard resulting tuples if they fail *qualifications*.
 - Delete attributes that are not in *target-list*.
 - If **DISTINCT** is specified, eliminate duplicate rows.
- ❖ This strategy is probably the least efficient way to compute a query! An optimizer will find more efficient strategies to compute *the same answers*.

Example Instances

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

- ❖ We will use these instances of the Sailors and Reserves relations in our examples.

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
28	yuppy	9	35.0
44	guppy	5	35.0

Using multiple tables

- ❖ Exercise: Find names of sailors who have reserved boat number 103.

Note 1: Sailor names only available at Sailors table

Note 2: We need reservation info as well...

Example using R1, S1 instances

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid AND R.bid=103
```

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/ 10/ 96
22	dustin	7	45.0	58	103	11/ 12/ 96
31	lubber	8	55.5	22	101	10/ 10/ 96
31	lubber	8	55.5	58	103	11/ 12/ 96
58	rusty	10	35.0	22	101	10/ 10/ 96
58	rusty	10	35.0	58	103	11/ 12/ 96

Note on Renaming

- ❖ Really needed only if the same relation appears twice in the FROM clause. The previous query can also be written as:

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid AND bid=103
```

OR

```
SELECT sname  
FROM   Sailors, Reserves  
WHERE  Sailors.sid=Reserves.sid  
       AND bid=103
```

*It is good style,
however, to use
range variables
always!*

Natural Join

- ❖ SQL:99 has a “natural join” clause:
“NATURAL” means equi-join for each pair of attributes with the same name

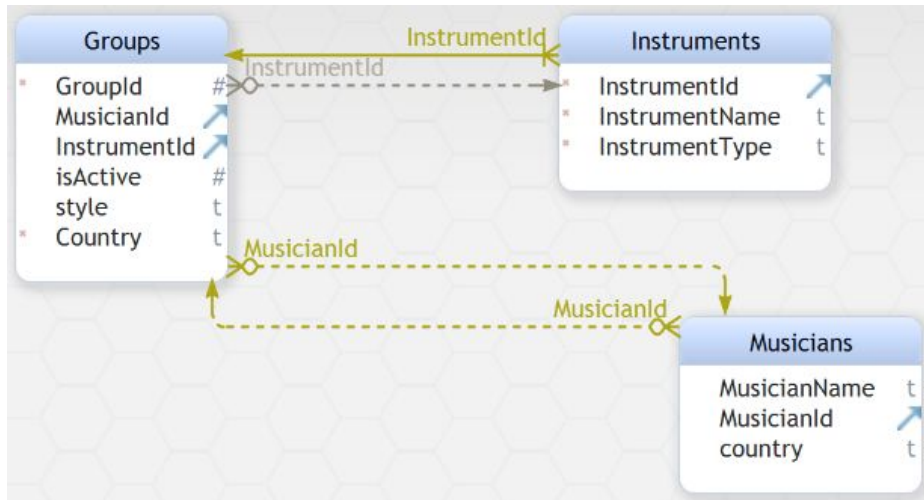
```
SELECT *  
FROM   Sailors  
natural join Reserves  
WHERE  bid=103
```

SID	SNAME	RATING	AGE	BID	DAY
22	dustin	7	45	103	08-OCT-98
31	lubber	8	55.5	103	06-NOV-98
58	rusty	10	35	103	12-NOV-98

Natural Join is evil

Sometimes you don't expect that 2 tables have 2 columns with the same name and that are different from primary key.

Example with this data model:



Natural Join is evil - Example

Musicians

MusicianId	MusicianNa	Age	Country
-----	-----	-----	-----
1	Jimmy	17	US
2	Paul	20	EN
3	Gustavo	60	BR

Groups

GroupId	MusicianId	Instr	style	Country
-----	-----	-----	-----	---
1	1	1	Blues	EN
2	3	2	Bossa Nova	BR
3	3	2	Bossa Nova	US

Question: I want to JOIN all Musicians with groups:

SELECT * FROM Groups JOIN Musicians WHERE Groups.MusicianId = Musicians.MusicianId

Groups

MusicianName	GroupId	Musicians.Country	Groups.Country
-----	-----	-----	-----
Jimmy	1	US	EN
Gustavo	2	BR	US
Gustavo	3	BR	BR



SELECT * FROM Groups NATURAL JOIN Musicians

Groups

MusicianName	GroupId	Musicians.Country	Groups.Country
-----	-----	-----	-----
Gustavo	3	BR	BR



❖ Exercise: Find sids of sailors who have reserved a red boat.

- Answer with manual Join:

```
SELECT  R.sid  
FROM    Boats B, Reserves R  
WHERE   B.bid = R.bid AND B.color = 'red';
```

- Answer with Natural Join:

```
SELECT  sid  
FROM    Boats NATURAL JOIN Reserves  
WHERE   color = 'red'
```

Expressions and Strings

- ❖ **AS** and **=** are two ways to name fields in result.
- ❖ **LIKE** is used for string matching. **`_`** stands for any one character and **`%`** stands for 0 or more arbitrary characters.

- ❖ Example:

```
SELECT S.age FROM Sailors S  
      WHERE S.sname LIKE 'B_%B'
```

Illustrates use of arithmetic expressions and string pattern matching: *Find records (age) for sailors whose names begin and end with B and contain at least three characters.*


Expressions and Strings (Example)

- ❖ Find sailors that have the character 'u' in their names

```
SELECT sname from Sailors  
where sname like '%u%'
```

SNAME
dustin
dustin
lubber
rusty
Brutus

Union (Addition)



- ❖ Goal: Find sid's of sailors who've reserved a red or a green boat
- ❖ **UNION**: Can be used to compute the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).
- ❖ If we replace **OR** by **AND** in the first version, what do we get?
- ❖ Also available: **EXCEPT** (What do we get if we replace **UNION** by **EXCEPT**?)

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND (B.color='red' OR B.color='green')
```

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND B.color='red'
```

```
UNION
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND B.color='green'
```

Union (Addition) (2)

- ❖ Find sids of sailors who have rating of 10 or reserved boat 104.


One attempt:

```
SELECT S.sid FROM Sailors S, Reserves R
WHERE S.sid = R.sid and (S.rating = 10 or R.bid = 104)
```

- with UNION:

```
SELECT  S.sid  FROM Sailors S
      WHERE  S.rating = 10
UNION
SELECT  R.sid  FROM Reserves R
      WHERE  R.bid = 104;
```

Except / minus




- ❖ To subtract a set from another, use:
 - EXCEPT : PostgreSQL
 - MINUS : MySQL et Oracle
- ❖ Find sids of sailors who have reserved green boats but not red boats.

```
SELECT  R.sid
        FROM    Reserves R, Boats B
        WHERE   R.bid = B.bid AND B.color = 'green'
MINUS
SELECT  R2.sid
        FROM    Boats B2, Reserves R2
        WHERE   R2.bid = B2.bid AND B2.color = 'red';
```

Note: don't forget to indent for lisibility

Intersect



- ❖ Find sid's of sailors who've reserved a red and a green boat
- ❖ **INTERSECT**: Can be used to compute the intersection of any two *union-compatible* sets of tuples.
- ❖ Included in the SQL/92 standard, but some systems don't support it.

```
SELECT S.sid  
FROM Sailors S, Boats B1, Reserves R1  
WHERE S.sid=R1.sid AND R1.bid=B1.bid  
AND (B1.color='red' AND B1.color='green')
```

```
SELECT S.sid  
FROM Sailors S, Boats B, Reserves R  
WHERE S.sid=R.sid AND R.bid=B.bid  
AND B.color='red'  
INTERSECT  
SELECT S.sid  
FROM Sailors S, Boats B, Reserves R  
WHERE S.sid=R.sid AND R.bid=B.bid  
AND B.color='green'
```

Nested Queries

- ❖ *Find names of sailors who've reserved boat #103:*

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN (SELECT R.sid
                FROM Reserves R
                WHERE R.bid=103)
```

```
SELECT S.sname
FROM Sailors S,
Reserves R
WHERE S.sid = R.sID
AND R.bid=103)
```

- ❖ A very powerful feature of SQL: a WHERE clause can itself contain an SQL query
- ❖ To find sailors who've *not* reserved #103, use NOT IN.
- ❖ To understand semantics of nested queries, think of a nested loops evaluation: *For each Sailors tuple, check the qualification by computing the subquery.*

Nested Queries with Correlation

Find names of sailors who've reserved boat #103:

```
SELECT S.sname  
FROM Sailors S  
WHERE EXISTS (SELECT *  
              FROM Reserves R  
              WHERE R.bid=103 AND S.sid=R.sid)
```



- ❖ **EXISTS** is another set comparison operator, like **IN**.
 - **IN**: Returns true if a specified value matches any value in a subquery or a list.
 - **Exists**: Returns true if a subquery contains any rows.

More on Set-Comparison Operators

- ❖ We've already seen IN, EXISTS. Can also use NOT IN, NOT EXISTS.
- ❖ Also available: *op* ANY, *op* ALL, *op* IN >, <, =, ≥, ≤, ≠
- ❖ Find sailors whose rating is greater than any ratings of some sailor called Horatio:

```
SELECT *  
FROM Sailors S  
WHERE S.rating > ANY (SELECT S2.rating  
                      FROM Sailors S2  
                      WHERE S2.sname='Horatio')
```

Rewriting *INTERSECT* Queries Using *IN*

Find sid's of sailors who've reserved both a red and a green boat:

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
      AND S.sid IN (SELECT S2.sid
                    FROM Sailors S2, Boats B2, Reserves R2
                    WHERE S2.sid=R2.sid AND R2.bid=B2.bid
                      AND B2.color='green')
```

- ❖ Similarly, *EXCEPT* queries re-written using *NOT IN*.
- ❖ To find *names* (not *sid's*) of Sailors who've reserved both red and green boats, just replace *S.sid* by *S.sname* in *SELECT* clause. (What about *INTERSECT* query? See next slide)

Intersect: Replacing sid by sname (1)

❖ The query will like this:

```
SELECT S.sname
  FROM Sailors S, Boats B, Reserves R
 WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
INTERSECT
SELECT S.sname
  FROM Sailors S, Boats B, Reserves R
 WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='green'
```

Does it work ?

Intersect: Replacing sid by sname (2)

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Figure 5.1 An Instance 53 of Sailors

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

Figure 5.2 An Instance *R2* of Reserves

red
green

<i>bid</i>	<i>bname</i>	<i>color</i>
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Figure 5.3 An Instance *B1* of Boats

Bug: Horatio will be in the result and shouldn't be.

Intersect: Replacing sid by sname (3)

To resolve earlier version that used INTERSECT and had small bug:

```
SELECT    S.sname
FROM  Sailors S WHERE    S.sid IN (
    (SELECT  R.sid
      FROM Boats B, Reserves R
      WHERE  R.bid = B.bid AND B.color = 'red')
  INTERSECT
  (SELECT  R2.sid
      FROM Boats B2, Reserves R2
      WHERE  R2.bid = B2.bid AND B2.color = 'green'));
```

Aggregate Operators

- ❖ Significant extension of relational algebra.

COUNT (*)
COUNT ([DISTINCT] A)
SUM ([DISTINCT] A)
AVG ([DISTINCT] A)
MAX (A)
MIN (A)

*single
column*

- ❖ Example: Minimum rating (worst) of the sailors:

```
SELECT MIN (S.rating)
FROM Sailors S
```

Aggregate Operators - Example

- ❖ Example: Find average age of sailors with a rating of 10.

```
SELECT  AVG (S.age)
        FROM    Sailors S
        WHERE   S.rating = 10;
```

- ❖ Example: Count the number of sailors.

```
SELECT  COUNT(*)
        FROM Sailors;
```

- ❖ Example: Count number of different sailors' names.

```
SELECT  COUNT(DISTINCT S.sname)
        FROM Sailors S;
```


Aggregate Operators - Example (2)

- ❖ Example: Find name and age of the oldest sailor(s)
- ❖ The first query has a bug!
(We'll look into the reason a bit later, when we discuss **GROUP BY**.)
- ❖ The third query is equivalent to the second query, and is allowed in the SQL/92 standard, but is not supported in some systems.

```
SELECT S.sname, MAX (S.age)
FROM Sailors S
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE S.age =
      (SELECT MAX (S2.age)
       FROM Sailors S2)
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE (SELECT MAX (S2.age)
      FROM Sailors S2)
      = S.age
```

Aggregate Operators - Example (3)

❖ Example: Find names of sailors who are older than oldest sailor with rating of 10.

❖ Solution using MAX:

```
SELECT S.sname FROM Sailors S
      WHERE S.age > (SELECT  MAX(S2.age)
                     FROM Sailors S2 WHERE  S2.rating = 10);
```

❖ Solution using ALL:

```
SELECT S.sname FROM Sailors S
      WHERE S.age > ALL (
                     SELECT  S2.age
                     FROM Sailors S2
                     WHERE  S2.rating = 10);
```

Queries With *GROUP BY* and *HAVING*

```
SELECT    [DISTINCT] target-list
FROM      relation-list
WHERE     qualification
GROUP BY  grouping-list
HAVING    group-qualification
```

- ❖ The *target-list* contains (i) attribute names (ii) terms with aggregate operations (e.g., MIN (S.age)).
 - The attribute list (i) must be a subset of *grouping-list*.
Intuitively, each answer tuple corresponds to a *group*, and these attributes must have a single value per group. (A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.)

Conceptual Evaluation

- ❖ The cross-product of *relation-list* is computed, tuples that fail *qualification* are discarded, 'unnecessary' fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
- ❖ The *group-qualification* is then applied to eliminate some groups. Expressions in *group-qualification* must have a *single value per group*!
- ❖ One answer tuple is generated per qualifying group.

For each red boat, find the number of reservations for this boat

```
SELECT B.bid, COUNT (*) AS scount  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```

Aggregate Operators - Example (4)

- ❖ Select * from sailors
order by rating

SID	SNAME	RATING	AGE
10	dustin	1	45
29	Brutus	1	33
85	Art	3	25.5
95	Bob	3	63.3
96	Frodo	3	25.5
22	dustin	7	45
64	Horatio	7	35
31	lubber	8	55.5
32	Andy	8	25.5
74	Horatio	9	35
71	Zorba	10	16
58	rusty	10	35

- ❖ Example: Number of all sailors
select count(*) from sailors → 12

- ❖ select rating, count(*) as count from
sailors **group by rating** order by rating

RATING	COUNT
1	2
3	3
7	2
8	2
9	1
10	2

Aggregate Operators - Example (4)

- ❖ Select * from sailors order by rating

SID	SNAME	RATING	AGE
10	dustin	1	45
29	Brutus	1	33
85	Art	3	25.5
95	Bob	3	63.3
96	Frodo	3	25.5
22	dustin	7	45
64	Horatio	7	35
31	lubber	8	55.5
32	Andy	8	25.5
74	Horatio	9	35
71	Zorba	10	16
58	rusty	10	35

- ❖ If we want to add sname to results we try this but it won't work:

select sname, rating, count(*) from sailors group by rating order by rating

SNAME	RATING	COUNT
Dustin, Brutus ??	1	2
Art, Bob, Frodo ??	3	3
Dustin, Horatio??	7	2
Lubber, Andy ??	8	2
Horatio	9	1
Zorba, rusty ??	10	2

Column List is Restricted

- ❖ Only one value is allowed per attribute, therefore the columns must either be in the group by list, or summarization (aggregate) functions like sum, count, avg, min, max
- ❖ sname cannot be put into column list
- ❖ Some database will accept it but will display only the first sname in the result table.

Aggregate Operators - Example (5)

❖ Select * from reserves order by bid =>

❖ Example: Find the number of past reservations for each boat ??

```
SELECT bid, COUNT(*)  
FROM reserves GROUP BY bid
```

BID	COUNT(*)
102	3
101	2
104	2
103	3

SID	BID	DAY
64	101	05-SEP-98
22	101	10-OCT-98
22	102	10-OCT-98
31	102	10-NOV-98
64	102	08-SEP-98
31	103	06-NOV-98
22	103	08-OCT-98
74	103	08-SEP-98
31	104	12-NOV-98
22	104	07-OCT-98

❖ Can we list boat names in the list as well?

```
SELECT r.bid, b.bname, count(*)  
FROM reserves r,boats b  
WHERE r.bid = b.bid  
GROUP BY r.bid, b.bname
```

Aggregate Operators - Example (6)

- ❖ List the ratings > 5 for which the average age is less than 40:

```
SELECT rating, AVG(age), COUNT(*)  
FROM sailors  
WHERE rating > 5  
GROUP BY rating  
HAVING AVG(age) < 40  
ORDER BY rating
```

RATING	AVG(AGE)	COUNT(*)
9	35	1
10	25.5	2

Let's detail this query in next slides

Aggregate Operators - Example (6)

❖ Step 1: SELECT rating, age FROM sailors

RATING	AGE
1	45
1	33
3	25.5
3	63.3
3	25.5
7	45
7	35
8	55.5
8	25.5
9	35
10	16
10	35

❖ Step 2: WHERE rating > 5

RATING	AGE
7	45
7	35
8	55.5
8	25.5
9	35
10	16
10	35

❖ Step 3: GROUP BY rating (it apply aggregators: avg and count)

RATING	AVG(AGE)	COUNT(
7	$(45+35)/2 = 40$	2
8	$(55.5+25.5)/2=40.5$	2
9	35	1
10	$(16+35)/2=25.5$	2

Aggregate Operators - Example (6)

❖ Step 4: $\text{HAVING AVG(AGE)} < 40$

RATING	AVG(AGE)	COUNT(
7	40	2
8	40.5	2
9	35	1
10	33.5	2

❖ Final result is:

RATING	AVG(AGE)	COUNT(*)
9	35	1
10	25.5	2

Find age of the youngest sailor with at least age 18, for each rating with at least 2 such sailors

```
SELECT S.rating, MIN (S.age)
      AS minage
FROM   Sailors S
WHERE  S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

*Answer
relation:*


rating	minage
3	25.5
7	35.0
8	25.5

Sailors instance:


<u>sid</u>	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5

Find age of the youngest sailor with age 18, for each rating with at least 2 such sailors.

rating	age
7	45.0
1	33.0
8	55.5
8	25.5
10	35.0
7	35.0
10	16.0
9	35.0
3	25.5
3	63.5
3	25.5



rating	age
1	33.0
3	25.5
3	63.5
3	25.5
7	45.0
7	35.0
8	55.5
8	25.5
9	35.0
10	35.0



rating	minage
3	25.5
7	35.0
8	25.5

Find age of the youngest sailor with age 18, for each rating with at least 2 sailors between 18 and 60.

```
SELECT S.rating, MIN (S.age)
      AS minage
FROM   Sailors S
WHERE  S.age >= 18 AND S.age <= 60
GROUP BY S.rating
HAVING COUNT (*) > 1
```

Sailors instance:

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5

*Answer
relation:*

rating	minage
3	25.5
7	35.0
8	25.5

*Find age of the youngest sailor with age > 18,
for each rating with at least 2 sailors (of any age)*

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age > 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT (*)
             FROM Sailors S2
             WHERE S.rating=S2.rating)
```

- ❖ Shows HAVING clause can also contain a subquery.
- ❖ What if HAVING clause is replaced by:
 - HAVING COUNT(*) >1

Summary on SET operators

- ❖ UNION: addition
- ❖ INTERSECT: subtraction
- ❖ CROSS JOIN: Multiplication

- ❖ Where is the Division ?
 - It doesn't exist in SQL.
 - But First what is the definition ?

Division in SQL: Definition

The general idea is the division of the dividend from a table which is divided with a divider table to obtain a quotient (a table result). The result is calculated from the values of a column to which the second column of the table dividend have all the values of the divider.

Reserves

sailor_name		boat_id
Peter	5	pedalo
Peter		Titanic
Peter		catamaran
Peter		jet-ski
Peter		pedalo
Anna	4	jet-ski
Anna		catamaran
Anna		catamaran
Anna		Yatch
Bob	1	Yatch

Boats

boat_id
jet-ski
Yatch
catamaran

```
SELECT sname
FROM Reserves R JOIN Boat B ON R.bid = B.bid
WHERE boat_id IN (SELECT bid FROM boats)
GROUP BY sailor_name
HAVING
    COUNT(DISTINCT(bid))
    =
    (SELECT count(*) FROM Boats)
```

Goal: Find sailors who've reserved all boats:

Division in SQL : 1st Method : Count

For sailors example:

```
SELECT sailor_name FROM Reserves  
      WHERE boat_id IN (SELECT boat_id FROM Boats)  
GROUP BY sailor_name  
HAVING COUNT(DISTINCT(boat_id)) = (SELECT COUNT(boat_id) FROM Boats)
```

Explanation:

Gets all sailors, count the number of different boats they have sailed and check if this number is equal to the number total of boat to sail.

Generic Formula:

```
SELECT group_id FROM Dividend  
      WHERE item_name IN (SELECT item_name FROM Divisor)  
GROUP BY group_id  
HAVING COUNT(DISTINCT(item_name)) = (SELECT COUNT(item_name) FROM Divisor)
```

Division in SQL : 2nd Method : Double Negation

The Query for sailors database:

```
SELECT DISTINCT(sailor_name) FROM Sailors AS S1
WHERE NOT EXISTS
    (SELECT * FROM Boats AS B
      WHERE NOT EXISTS
        (SELECT * FROM RESERVES AS R
         WHERE R.sid=S1.sid AND R.bid = B.bid)
    )
```

Explanation:

Each sailor look at the boats and say:

There is no boat here

That I haven't sail/reserved

Null Values

- ❖ Field values in a tuple are sometimes *unknown* (e.g., a rating has not been assigned) or *inapplicable* (e.g., no spouse's name).
 - SQL provides a special value *null* for such situations.
- ❖ The presence of *null* complicates many issues. E.g.:
 - Special operators needed to check if value is/is not *null*.
 - Is *rating* > 8 true or false when *rating* is equal to *null*? What about *AND*, *OR* and *NOT* connectives? ($5 < \text{null} \Rightarrow \text{unknown}$)
 - We need a 3-valued logic (true, false and *unknown*).
 - Meaning of constructs must be defined carefully. (e.g., WHERE clause eliminates rows that don't evaluate to true.)
 - New operators (in particular, *outer joins*) possible/needed.

Three-Valued Logic

OR	Unknown	True	False
Unknown	Unknown	True	Unknown
True	True	True	True
False	Unknown	True	False

AND	Unknown	True	False
Unknown	Unknown	Unknown	False
True	Unknown	True	False
False	False	False	False

SQL and NULL Values

WHERE clause is unknown? SQL query eliminates those rows (it evaluates to false)

DISTINCT - two rows are duplicates if corresponding columns are equal or both are null

SQL and NULL Values (continued)

Arithmetic operators - return NULL if argument is null

COUNT(*) - includes null values

All other aggregate operators discard null values.

Integrity Constraints (Review)

- ❖ An IC describes conditions that every *legal instance* of a relation must satisfy.
 - Inserts/deletes/updates that violate IC's are disallowed.
 - Can be used to ensure application semantics (e.g., *sid* is a key), or prevent inconsistencies (e.g., *sname* has to be a string, *age* must be < 200)
- ❖ Types of IC's: Domain constraints, primary key constraints, foreign key constraints, general constraints.
 - *Domain constraints*: Field values must be of right type. Always enforced.

General Constraints

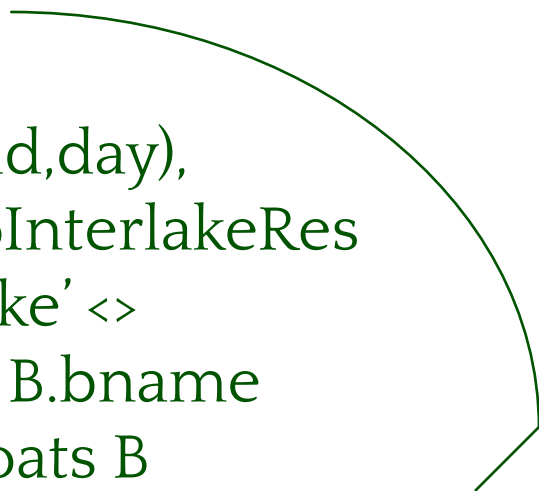
- ❖ Useful when more general ICs than keys are involved.
- ❖ Can use queries to express constraint.
- ❖ Constraints can be named.

```
CREATE TABLE Sailors
```

```
( sid INTEGER,  
  sname CHAR(10),  
  rating INTEGER,  
  age REAL,  
  PRIMARY KEY (sid),  
  CHECK ( rating >= 1  
         AND rating <= 10 )
```

```
CREATE TABLE Reserves
```

```
( sname CHAR(10),  
  bid INTEGER,  
  day DATE,  
  PRIMARY KEY (bid,day),  
  CONSTRAINT noInterlakeRes  
  CHECK ( `Interlake' <>  
        ( SELECT B.bname  
          FROM Boats B  
          WHERE B.bid=bid)))
```



JOINS

```
SELECT (column_list)
FROM table_name
  [INNER | {LEFT | RIGHT | FULL } OUTER] JOIN table_name
    ON qualification_list
WHERE ...
```

Explicit join semantics needed unless it is an INNER
join

(INNER is default)

Inner Join

Only the rows that match the search conditions are returned.

```
SELECT s.sid, s.name, r.bid
```

```
FROM Sailors s INNER JOIN Reserves r
```

```
ON s.sid = r.sid
```

Returns only those sailors who have reserved boats

SQL-92 also allows:

```
SELECT s.sid, s.name, r.bid
```

```
FROM Sailors s NATURAL JOIN Reserves r
```

Inner Join : Example

SELECT s.sid, s.name, r.bid

FROM Sailors s INNER JOIN Reserves r

ON s.sid = r.sid

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

Result:

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103

Left Outer Join

Left Outer Join returns all matched rows, plus all unmatched rows from the table on the left of the join clause.

(use nulls in fields of non-matching tuples)

```
SELECT s.sid, s.name, r.bid  
FROM Sailors s LEFT OUTER JOIN Reserves r  
ON s.sid = r.sid
```

Returns all sailors & information on whether they have reserved boats

Left Outer Join : Example

SELECT s.sid, s.name, r.bid

FROM Sailors s LEFT OUTER JOIN Reserves r

ON s.sid = r.sid

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

Result:

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103
31	Lubber	NULL

Right Outer Join

Right Outer Join returns all matched rows, plus all unmatched rows from the table on the right of the join clause

```
SELECT r.sid, b.bid, b.name  
FROM Reserves r RIGHT OUTER JOIN Boats b  
ON r.bid = b.bid
```

Returns all boats & information on which ones are reserved.

Right Outer Join: Example

SELECT r.sid, b.bid, b.name

FROM Reserves r RIGHT OUTER JOIN Boats b

ON r.bid = b.bid

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Result:

r.sid	b.bid	b.name
22	101	Interlake
NULL	102	Interlake
95	103	Clipper
NULL	104	Marine

Full Outer Join

Full Outer Join returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT r.sid, b.bid, b.name  
FROM Reserves r FULL OUTER JOIN Boats b  
ON r.bid = b.bid
```

Returns all boats & all sailors & all links between reservation if there are

Full Outer Join

Find sids of sailors who have rating of 10 or reserved boat 104

What about sailors who have not reserved any boats?

What if we have some old reservations in the database for which the Sailor info is missing?

Use full outer join that includes rows with no match:

```
SELECT S.sid
```

```
FROM Sailors S full outer join Reserves R on (S.sid = R.sid)
```

```
WHERE S.rating = 10 or R.bid = 104
```

Full Outer Join : Example

SELECT r.sid, b.bid, b.name

FROM Reserves r FULL OUTER JOIN Boats b

ON r.bid = b.bid

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
NULL	102	Interlake
95	103	Clipper
NULL	104	Marine


Result:

Note: in this case it is the same as the ROJ because bid is a foreign key in reserves, so all reservations must have a corresponding tuple in boats.

Full Outer Join : Example (2)

- Example will Null on left and right tables:

event_id	description	user_id
1	User Click	1
2	System Check	NULL



user_id	name
1	Charlie Brown
2	Snoopy

event_id	description	events.user_id	users.user_id	name
1	User Click	1	1	Charlie Brown
2	System Check	NULL	NULL	NULL
NULL	NULL	NULL	2	Snoopy

Summary

INNER JOIN	<div>1 2 3</div>	INNER JOIN	<div>A B C</div>	=	<div>1 B 2 A</div>	Only returns rows that meet the join condition
RIGHT OUTER JOIN	<div>1 2 3</div>	RIGHT OUTER JOIN	<div>A B C</div>	=	<div>1 B 2 A C</div>	Returns all rows from the table on the right side of JOIN and matched rows from the left side of the JOIN
LEFT OUTER JOIN	<div>1 2 3</div>	LEFT OUTER JOIN	<div>A B C</div>	=	<div>1 B 2 A 3</div>	Returns all rows from the table on the left side of JOIN and matched rows from the right side of the JOIN
FULL OUTER JOIN	<div>1 2 3</div>	FULL OUTER JOIN	<div>A B C</div>	=	<div>1 B 2 A 3 C</div>	Returns all rows from both sides even if join condition is not met
CROSS JOIN	<div>1 2 3</div>	CROSS JOIN	<div>A B C</div>	=	<div>1 A 1 B 1 C 2 A 2 B 2 C 3 A 3 B 3 C</div>	Cartesian product between the two sides is a join but without a join condition. Returns all rows joined from both sides

Advanced Summary

LEFT JOIN



Everything on the left
+
anything on the right that
matches

```
SELECT *  
FROM TABLE_1  
LEFT JOIN TABLE_2  
ON TABLE_1.KEY = TABLE_2.KEY
```

ANTI LEFT JOIN



Everything on the left
that is NOT on the right

```
SELECT *  
FROM TABLE_1  
LEFT JOIN TABLE_2  
ON TABLE_1.KEY = TABLE_2.KEY  
WHERE TABLE_2.KEY IS NULL
```

RIGHT JOIN



Everything on the right
+
anything on the left that matches

```
SELECT *  
FROM TABLE_1  
RIGHT JOIN TABLE_2  
ON TABLE_1.KEY = TABLE_2.KEY
```

ANTI RIGHT JOIN



Everything on the right
that is NOT on the left

```
SELECT *  
FROM TABLE_1  
RIGHT JOIN TABLE_2  
ON TABLE_1.KEY = TABLE_2.KEY  
WHERE TABLE_1.KEY IS NULL
```

OUTER JOIN



Everything on the right
+
Everything on the left

```
SELECT *  
FROM TABLE_1  
OUTER JOIN TABLE_2  
ON TABLE_1.KEY = TABLE_2.KEY
```

ANTI OUTER JOIN



Everything on the left and right
that is unique to each side

```
SELECT *  
FROM TABLE_1  
OUTER JOIN TABLE_2  
ON TABLE_1.KEY = TABLE_2.KEY  
WHERE TABLE_1.KEY IS NULL  
OR TABLE_2.KEY IS NULL
```

INNER JOIN



Only the things that match on the
left AND the right

```
SELECT *  
FROM TABLE_1  
INNER JOIN TABLE_2  
ON TABLE_1.KEY = TABLE_2.KEY
```

CROSS JOIN



All combination of rows from the
right and the left (cartesian
product)

```
SELECT *  
FROM TABLE_1  
CROSS JOIN TABLE_2
```

Data dictionary

- ❖ Tables in which descriptions of the objects of the base are stored.
 - It is held up to date automatically by the DBMS.
 - These tables can be consulted by SQL language.
- ❖ Examples:
 - `DICTIONARY (DICT)` : dictionary view
 - `USER_TABLES` : tables and views created by the current user
 - `USER_CATALOG (CAT)` : tables and views on which the user has rights, other than the tables and views of the dictionary of the data
 - `USER_TAB_COLUMNS (COLS)` : columns of the tables or views created by the user
 - `USER_TAB_PRIVS` : objects on which the user is donor or receiver of rights

Exercise

- ❖ We have a database which manage reservation places for theatrical performance (spectacles) .

Spectacles (NumSpec, NameSpec, DatBegSpec, DateEndSpec)

Performance (NumRep, NumSpec, DatPerf, HourPerf)

Place (NumRep, NumPlace, Price)

Book (NumRes, NumRep, NumPlace, NameDem, TelDem)

- ❖ Answer following requests:

- All performances of spectacle N°13 between 15 of september 2005 and 15 of January 2006 or between, 15 of Février 2006 and 15 of May 2006.
- All performances of the spectacle “Revisor” between 15 September 2005 and 15 January 2006.
- Find all persons that have booked a place for the last performance of 2005.
- Find all persons that have booked a place for the performance N°1 and N°2.

Exercise

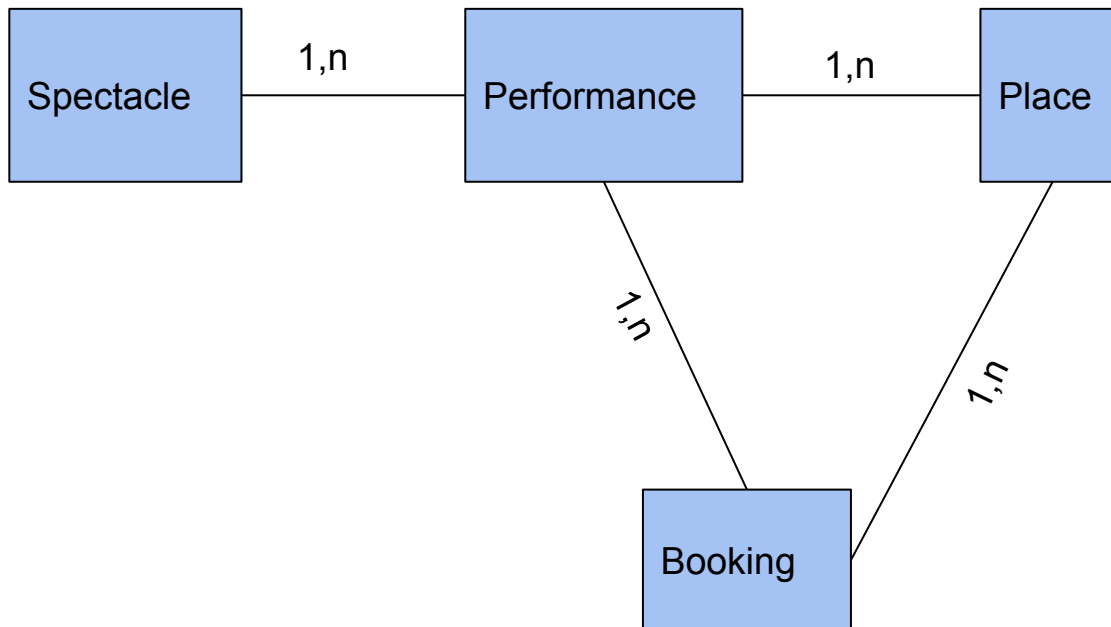
- ❖ We have a database which manage reservation places for theatrical performance (spectacles) .

Spectacles (NumSpec, NameSpec, DatBegSpec, DateEndSpec)


Performance (NumRep, NumSpec, DatPerf, HourPerf)

Place (NumRep, NumPlace, Price)

Booking (NumRes, NumRep, NumPlace, NameDem, TelDem)



- ChessPlayer (cid: Integer, cname: String, age: Integer, country: String)
- Match (cid: Integer, tid: Integer, isVictory: Boolean, round: String)
- Tournament (tid: Integer, tname: String, country: String, type: String)



- ❖ All performances of spectacle N°13 between 15 of september 2005 and 15 of January 2006 or between, 15 of Février 2006 and 15 of May 2006:

```
SELECT *  
FROM Performance  
WHERE DatPerf BETWEEN '15-SEP-2005' AND '15-JAN-2006'  
OR DatPerf BETWEEN '15-FEV-2006' AND '15-MAI-2006'  
AND NumSpec = 13;
```

- ❖ All performances of the spectacle “Revisor” between 15 September 2005 and 15 January 2006:

```
SELECT DatPerf, HeurPerf  
FROM Performance R, Spectacle S  
WHERE DatPerf BETWEEN '15-SEP-05' AND '15-JAN-06'  
AND NameSpec = 'Revisor'  
AND S.NumSpec = R.NumSpec ;
```

- ❖ Find all persons that have booked a place for the last performance of 2005:

```
SELECT NameDem
FROM Booked B, Performance P
WHERE P.NumRep = B.NumRep
AND P.DatPerf = ( SELECT MAX(DatPerf)
FROM Performance
WHERE DatPerf<='31-12-2005' );
```

- ❖ Find all persons that have booked a place for the performance N°1 and N°2:

```
SELECT DISTINCT NameDem, TelDem
FROM Performance
WHERE NumRep = 1
INTERSECT
SELECT DISTINCT NameDem, TelDem
FROM Performance
WHERE NumRep = 2 ;
```

- ❖ Find all persons that have booked a place for the last performance of 2005:

```
SELECT NameDem
FROM Booked B, Performance P
WHERE P.NumRep = B.NumRep
AND P.DatPerf = ( SELECT MAX(DatPerf)
FROM Performance
WHERE DatPerf<='31-12-2005' );
```

- ❖ Find all persons that have booked a place for the performance N°1 and N°2:

```
SELECT DISTINCT NameDem, TelDem
FROM Performance
WHERE NumRep = 1
INTERSECT
SELECT DISTINCT NameDem, TelDem
FROM Performance
WHERE NumRep = 2 ;
```

Games to learn SQL

Here are some games to learn/practice SQL:

<https://datalemur.com/blog/games-to-learn-sql>